

University of
Strathclyde
Glasgow

DEPARTMENT OF MATHEMATICS AND STATISTICS

Statistical and Machine Learning Approaches to Premier
League Betting:

A Multi-Market Optimization Framework

By

Christopher McCallan

202356850

MSc Applied Statistics
2024/25

Statement of work in project

The work contained in this project is that of the authors and where material from other sources has been incorporated full acknowledgement has been given.

Christopher McCallan

Signed:

Print Name: Christopher McCallan:

Date: 20/07/2025

Abstract.....	6
1. Introduction	6
1.1 Research Background and Motivation	6
1.2 Research Objectives and Questions	8
Key Research Questions:.....	8
Scope Inclusions:.....	8
Key Limitations:.....	9
2. Literature Review.....	9
2.1 Historical Development of Sports Statistics.....	9
2.2 Traditional Statistical Approaches in Football.....	9
2.2.1 Poisson Distribution Models	9
2.2.2 Maher	10
2.2.3 Dixon & Coles	11
2.3 Alternative Distributional Approaches.....	12
2.4 Modern Machine Learning Applications in Football Analytics	13
2.4.1 Covid-19 Impact on Football Analysis.....	14
2.4.2 Contemporary Betting Syndicates and Professional Approaches.....	14
2.4.3 Feature Engineering in Recent Literature	16
2.5 Research Focus and Expected Contributions.....	18
2.5.1 Building Upon Established Foundations	18
2.5.2 Contemporary Modeling Approach	19
2.5.3 Practical Implementation Focus.....	19
2.5.4 Multi-Market Optimization Strategy.....	19
2.5.5 Research Aims and Implementation Goals.....	20
Implementation Validation Criteria:.....	20
3 Theoretical Framework	21
3.1 Statistical Foundations for Sports Prediction	21
3.1.1 Probability Theory and Rare Events	21
3.1.2 Statistical Inference and Model Selection	21
3.1.3 Time Series Considerations.....	21
3.2 Machine Learning Theory and Feature Engineering	22
3.2.1 Bias-Variance Decomposition	22
3.2.2 Feature Engineering and Representation Learning	22
3.2.3 Ensemble Methods and Model Combination	22
3.3 Market Efficiency and Value Theory.....	23
3.3.1 Efficient Market Hypothesis in Sports Betting	23
3.3.2 Expected Value and Optimal Betting Strategy	24
3.3.3 Multi-Market Optimization Approach	25
3.3.4 Risk Management and Practical Implementation.....	25
4 Data and Methodology.....	26
4.1 Data Sources and Collection.....	26
4.1.1 Match Results and Betting Data	26
4.1.2 Advanced Performance Metrics.....	27
4.1.3 Data Integration Framework	29
4.2 Feature Engineering	30
4.2.1 Traditional Statistical Features	30

4.2.2 Contemporary Performance Metrics	30
<i>Expected Value Framework</i>	30
<i>Possession Quality Indicators</i>	31
<i>Defensive Intensity Metrics</i>	31
4.2.3 Newly Promoted Team Handling	31
<i>Championship-to-Premier League Translation</i>	31
<i>Implementation Protocol:</i>	31
4.2.4 Market Intelligence Features	31
<i>Price Movement Indicators</i>	32
<i>Cross-Market Arbitrage</i>	32
4.2.5 Feature Selection and Dimensionality Reduction	32
<i>Correlation Analysis and Multicollinearity Detection</i>	32
<i>Principal Component Analysis</i>	33
<i>Recursive Feature Elimination with Cross-Validation</i>	33
<i>Feature Importance Quantification</i>	33
<i>Temporal Stability Assessment</i>	34
<i>Market-Specific Feature Selection</i>	34
<i>Final Feature Selection Protocol</i>	34
4.3 Experimental Design	35
4.3.1 Model Development Framework	35
4.3.2 Model Validation Methodology	35
4.3.3 Covid-19 Impact Analysis	37
<i>Dual-Dataset Experimental Design:</i>	37
<i>Statistical Testing Framework</i>	37
4.3.4 Implementation Considerations	37
<i>Reproducibility Measures</i>	37
5 Model Implementation and Results	38
5.1 Implementation Overview	38
5.1.1 Development Environment	38
5.1.2 Data Pipeline and Temporal Splitting	38
5.1.3 Feature Engineering Strategy	39
5.2 Traditional Poisson Model Results	39
5.2.1 Model Training and Convergence	39
5.2.2 Team Parameter Analysis	40
<i>Table 5.1: Parameter Distribution Statistics</i>	40
5.2.3 Predictive Performance	40
<i>Table 5.2: Poisson Model Test Set Performance</i>	40
5.3 Dixon-Coles Enhanced Model Results	40
5.3.1 Model Parameters and Optimization	40
5.3.2 Time Weighting Effect	41
<i>Table 5.3: Match Weighting by Recency</i>	41
5.3.3 Performance Improvement	41
<i>Table 5.4: Dixon-Coles vs Poisson Comparison</i>	41
5.3.4 Statistical Significance Testing	42
5.4 Model C: Contemporary ML Ensemble Results	42
5.4.1 Training Configuration	42
5.4.2 Multi-Market Performance	42
<i>Table 5.5: Model C Performance Across Markets</i>	42

5.4.3 Feature Importance Analysis	43
5.5 Comparative Model Analysis.....	45
5.5.1 Overall Performance Comparison.....	45
Table 5.6: Comprehensive Model Comparison.....	45
5.5.2 Statistical Significance	46
Table 5.7: Extended Statistical Comparisons	46
Mean Difference	46
t- p- Cohen's statistic value d.....	46
Interpretation	46
5.6 Value Betting Analysis.....	47
5.6.1 Value Bet Identification	47
5.6.2 Betting Simulation Results.....	47
5.6.3 Monthly Performance Analysis	47
5.7 Covid Era Analysis.....	48
5.7.1 Home Advantage Erosion	48
Period Matches Home Win Rate Home PPG Away PPG Home Advantage.....	48
5.7.2 Statistical Distribution Shifts	48
5.7.3 Model Performance During COVID	49
Table 5.11: Model Degradation During Pandemic	49
5.7.4 Model Comparison: With vs Without Covid Data	49
5.8 Implementation Considerations	50
5.8.1 Computational Requirements	50
Table 5.12: Resource Utilization Comparison.....	50
5.8.2 Practical Deployment Challenges.....	50
5.9 Key Findings	51
5.9.1 Model Performance Hierarchy	51
5.9.2 Theoretical Contributions.....	51
5.9.3 Practical Insights.....	51
6 Discussion and Analysis.....	51
6.1 Theoretical Implications	52
6.1.1 The Obsolescence of Classical Assumptions	52
6.1.2 Market Efficiency and Information Integration.....	53
6.1.3 The Limits of Machine Learning in Structured Domains	53
6.2 Practical Performance Analysis	54
6.2.1 The Value Identification Paradox	54
6.2.2 Realistic Performance Expectations	54
6.2.3 Multi-Market Optimization Validation	54
6.3 COVID-19 as a Natural Experiment	55
6.3.1 The Fragility of Statistical Assumptions	55
6.3.2 Implications for Adaptive Modeling	55
6.4 Methodological Contributions and Limitations.....	55
6.4.1 Contributions to Sports Analytics Literature	55
6.4.2 Acknowledged Limitations	56
6.5 Implications for Stakeholders	56
6.5.1 For Academic Researchers	56
6.5.2 For Practitioners.....	57
6.5.3 For the Betting Industry	57

6.6 Future Research Directions.....57
6.6.1 Immediate Extensions 57
6.6.2 Fundamental Questions..... 57
6.7 Concluding Thoughts58
7 **Conclusions, Enhancements and Future Research..... 58**
7.1 Summary of Findings..... 58
7.2 Contributions to Knowledge 59
Theoretical Contributions: 59
Methodological Contributions: 59
7.3 Practical Recommendations 59
7.4 Future Research Directions 60
Immediate Extensions: 60
Longer-Term Opportunities: 60
7.5 Final Reflections 60
Appendix 60
References 61

Abstract

This dissertation investigates the application of statistical modeling techniques to predict Premier League football match outcomes and identify profitable betting opportunities across multiple betting markets. The research systematically compares three distinct approaches: the traditional Poisson distribution model, the enhanced Dixon-Coles framework, and Model C - a novel machine learning ensemble approach that integrates advanced performance metrics with multi-market optimization strategies.

The study utilizes two Premier League datasets, one spanning from the 2019/20 to 2024/25 season and another beginning from 2020/21, to develop and validate predictive models. The rationale for testing models with different starting seasons is to explore how the anomalous COVID-19 season affects statistical modeling, given how match conditions changed drastically in an abnormal manner. Data sources include historical match results from football-data.co.uk and comprehensive performance metrics from fbref.com, encompassing over 2280 matches and over 120+ statistical features per team-season.

Key contributions include: (1) development of a contemporary feature engineering pipeline leveraging advanced metrics like expected goals (xG) and tactical style indicators; (2) a systematic framework for handling newly promoted teams using Championship performance scaling; (3) implementation of multi-market optimization that extends beyond traditional match outcome prediction to identify value across over/under goals and both teams to score markets; and (4) preliminary analysis of COVID-19's impact on model performance. The research addresses practical implementation challenges while maintaining academic rigor, providing insights valuable to both sports analytics researchers and practitioners interested in the intersection of machine learning and sports betting markets.

Model C achieved 53.2% accuracy on match outcomes and identified value betting opportunities yielding 146.8% ROI in back-testing, though practical constraints including account restrictions significantly impact real-world implementation

1. Introduction

1.1 Research Background and Motivation

The book *Moneyball: The Art of Winning an Unfair Game* released in 2003 by author Michael Lewis, popularised the heralding of a new era in sports. Detailing Major League Baseball's Oakland Athletics and General Manager Billie Beane's attempt to assemble a World Series winning team on a shoestring budget, the book's release and subsequent Hollywood film depicts how Beane used sabermetrics (James 1980) and statistical analysis over traditional baseball knowledge and scouting to find players capable of beating the most valuable and best paid teams in baseball at a fraction of the price their statistics would suggest they were worth.

The Oakland A's would make the playoffs in 2002 and 2003 on a salary budget of \$44m (compared to the New York Yankees \$125m budget) with Beane being the face of a new era of data analysis in baseball.

Beane's influence didn't stop at baseball however. After the release of the book he met with leading figures in football. Arsene Wenger and Sir Alex Ferguson are both known admirers of Beane's work. John C. Henry, owner of the Boston Red Sox and later Liverpool Football Club, attempted to hire Beane to the Red Sox. Liverpool later implemented a statistical and data science led structure that has major influence on player recruitment, managerial appointments and day to day tactics and training.

It would however be unfair to claim football had never sought to employ data driven models before Moneyball's release. Sam Allardyce, whose managerial career in England spans 12 clubs and briefly the England national team, has long been a proponent of data driven tactics and player recruitment (despite being unfairly maligned as a long ball merchant), using statistical analysis going so far back as the late 90s.

Given the popularity of football, particularly since the Sky deal bringing in the era of the Premier League in the early 90s, it was inevitable that other industries would seek to use data as means to their own ends. Sports betting has been popular in the UK for centuries, with horse racing being popular since the 1500s. The Gaming Act of 1845 formalised rules and criminalised cheating, and in the early 20th century football pools were a popular vice among the working class as Association Football's popularity exploded.

In the UK, sports betting is predicted to be worth over £4 billion annually in the next few years (Williams 2024). Despite growing concerns over societal issues such as gambling addiction and financial implications the industry's revenues are increasing year on year. Bet365's owners, the Coates family based in Stoke, are regularly in the UK's top 5 tax payers lists due to their huge revenues and profits.

While the industry grows in the UK and globally (some US states have recently relaxed or decriminalised some forms of sports betting, having historically employed a blanket ban on the industry), it's clear that it isn't going away. So the question becomes, is it possible to beat the bookmakers at their own game? In their UK, the weekend accumulator is a common popular Saturday activity. Largely based on high level statistics such as team form, as well as a gut instinct, punters hope for a big payday betting leagues across the UK and Europe.

While the vast majority of these lose, there are some serial big winners on football betting. Brentford FC's owner Matthew Benham famously made his fortune betting on football through his company SmartOdds. Likewise, Benham's former colleague and now bitter rival Tony Bloom, owner of Brighton & Hove Albion FC and betting company Starlizard, made his fortune in a similar manner. Both employ statistical models to predict results of football matches across the globe, as well as driving player recruitment at their respective clubs. Clearly, the opportunities are there to make

money in the sports betting markets, it's seemingly just a case of having enough data and mathematical prowess (and doubtlessly some luck) to do so.

This research aims to identify market inefficiencies in Premier League betting markets and evaluate whether sophisticated statistical modeling can generate sustainable profits. This work contributes to the growing field of sports analytics by providing a comprehensive comparison of modeling approaches and demonstrating the practical application of advanced statistical techniques in a real-world betting context.

1.2 Research Objectives and Questions

This project aims to utilize the wealth of data widely available to bridge the gap between theoretical statistical modeling and practical application in football prediction and betting market analysis. In doing so we evaluate three distinct modeling approaches:

1. **Traditional Poisson Model:** The foundational approach assuming independent goal distributions
2. **Dixon-Coles Enhanced Model:** Incorporating time-decay and low-score adjustments
3. **Model C - Contemporary Multi-Market Optimization:** Our novel framework that integrates:
 - Advanced performance metrics (xG, xA, possession quality metrics)
 - Machine learning ensemble methods (XGBoost + Neural Networks)
 - Multi-market value identification across multiple betting products
 - Dynamic feature selection based on market type
 - Bayesian confidence intervals for risk-adjusted betting

Model C represents the primary methodological contribution, designed specifically for practical betting implementation rather than academic accuracy metrics alone.

Key Research Questions:

1. Can Model C achieve a statistically significant improvement in log-loss over Dixon-Coles when applied to contemporary Premier League data?
2. Which betting markets show greatest inefficiency as measured by positive expected value?
3. Does excluding COVID-affected data improve model calibration by a statistically significant margin?
4. Can multi-market optimization strategies outperform traditional single-market approaches in terms of risk-adjusted returns?

Scope Inclusions:

- Pre-match betting markets only
- Statistical modeling approaches (Poisson, Dixon-Coles, Model C)
- Multiple data sources combining results with advanced performance metrics

- Practical betting strategy implementation and back-testing
- Newly promoted team handling methodology

Key Limitations:

- Premier League only (2019/20 - 2024/25 seasons)
- Exclusion of in-play betting opportunities
- Limited consideration of external factors such as injuries, transfers and managerial changes
- Sample size constraints (~2,280 matches) requiring careful statistical treatment

2. Literature Review

2.1 Historical Development of Sports Statistics

The application of statistical methods to sports analysis has deep historical roots, extending back to Henry Chadwick's development of the box score in baseball during the 1850s. However, the systematic application of probability theory to football prediction began considerably later, with Charles Reep's pioneering work in the 1950s providing the first comprehensive statistical analysis of association football.

Reep's methodology, focusing on possession statistics and pass completion rates, established fundamental principles that continue to influence modern analytics. His observation that most goals originate from moves of three passes or fewer challenged traditional tactical thinking and demonstrated the potential for statistical insight to contradict conventional wisdom (Reep & Benjamin, 1968).

The evolution from descriptive to predictive analytics gained momentum with the work of Moroney (1956), who first applied the Poisson distribution to football scoring patterns. This seminal contribution established the theoretical foundation for treating goals as independent random events, despite the obvious limitations of this assumption in real-world scenarios.

2.2 Traditional Statistical Approaches in Football

2.2.1 Poisson Distribution Models

The Poisson distribution is frequently used to model the number of goals scored in a football match. It assumes that goals occur independently with a constant average rate. With this framework, the probability of a team scoring x goals is given by:

$$P(X = x) = \frac{e^{-\lambda} \lambda^x}{x!}$$

This assumes that goals occur independently at a constant average rate λ . Early work by **Moroney (1951)** and **Maher (1982)** laid the groundwork for using Poisson models in match outcome prediction. These models calculate match scorelines such as 1–0 or 2–2 and are widely used for computing probabilities in betting markets.

Model	Key Innovation	Performance Gain	Practical Limitation	Section
	Attack/defence parameters,	+12% vs goals average	Assumes independence	2.2.2
Maher (1982)	home advantage			
Dixon & Coles (1997)	Time-weighting, draw inflation	+8% vs Maher	Complex parameter estimation	2.2.3
Karlis & Ntzoufras (2003)	Bivariate Poisson, score correlation	+3% vs D&C	Computationally intensive	2.3
Baio & Blangiardo (2010)	Bayesian hierarchical, uncertainty	Similar accuracy, better intervals	Slow inference	2.3
Model (2024)	C ML ensemble + multi-market	+8.3% vs baseline	Requires extensive features	5.4

Table 1

2.2.2 Maher

Maher's (1982) comprehensive framework extended Moroney's work by introducing separate attack and defence strength parameters, enabling more sophisticated team comparisons and prediction capabilities. The basic Poisson model framework assumes that the number of goals scored by each team follows independent Poisson distributions with parameters determined by the attacking strength of the scoring team and the defensive strength of the opposing team. Mathematically, this is expressed as:

Goals_{Home} ~ Poisson($\alpha_{\text{home}} \times \beta_{\text{away}} \times \gamma$)

Goals_{Away} ~ Poisson($\alpha_{\text{away}} \times \beta_{\text{home}}$)

where α represents attack strength, β represents defensive strength, and γ represents home advantage.

This approach, modeling home and away goal distributions as independent Poisson processes, became the standard methodology for football prediction models throughout the 1980s and 1990s.

While Maher's (1982) model marked a significant improvement over the basic Poisson framework by introducing team-specific attack and defence parameters, it suffers from several key limitations. Most notably, it assumes that goals scored by the home and away teams are **independent**, an assumption that fails to capture the interactive and strategic nature of football matches. For instance, a team taking the lead may adopt a more defensive posture, which influences the goal-scoring probability of both teams — an effect not captured in Maher's model.

Additionally, the model **treats all matches equally**, ignoring the fact that team performance can change significantly over time due to injuries, form, or tactical adjustments. As such, older matches may distort parameter estimates if not downweighted appropriately. Finally, Maher's model **underpredicts the frequency of lowscoring draws**, particularly 0–0 and 1–1 scorelines, which occur more frequently in practice than the independent Poisson assumption allows. These shortcomings limit its predictive accuracy in real-world football betting applications.

2.2.3 Dixon & Coles

While Maher's (1982) model introduced team-specific attack and defence parameters and a home advantage multiplier, it retained two key assumptions: that goals scored by the two teams are independent and that all matches contribute equally to parameter estimation regardless of how long ago they were played. The **Dixon & Coles (1997)** model extended Maher's framework to address these limitations, leading to improved prediction accuracy in football score modelling and betting market evaluation.

Improvement 1: Time Decay

Football teams evolve over time due to transfers, injuries, managerial changes, and form. Dixon & Coles introduced a time-weighted likelihood, down-weighting older matches so that recent results have more influence on the parameter estimates.

A decay factor is introduced such that the contribution of a match at time t is weighted by:

$$w_t = \exp(-\phi(T - t))$$

Where:

- T : time of the most recent match
- t : time of the historical match
- ϕ = decay parameter controlling the rate of historical discounting

This ensures that newer matches are more influential, allowing the model to adapt more quickly to changes in team performance.

Improvement 2: Scoreline Adjustment for Low-Scoring Draws

The original Maher model assumes that the number of goals scored by the home and away teams are independent. However, in reality, draws—particularly 0–0 and 1–1—occur more frequently than predicted by the independent Poisson assumption.

Dixon & Coles introduced a scoreline adjustment term, $\tau(x,y)$ to modify the likelihood when both teams score 0 or 1 goals. The joint probability of a scoreline is

$$P(X = x, Y = y) = \begin{cases} \tau(x, y) \cdot P_{\text{ind}}(x, y) & \text{if } x, y \in \{0, 1\} \\ P_{\text{ind}}(x, y) & \text{otherwise} \end{cases}$$

The Dixon & Coles model represents a pivotal advancement in football score prediction, addressing key shortcomings in Maher's earlier framework. By introducing a time decay factor, it allows the model to dynamically adjust to evolving team form, giving more weight to recent performances. Moreover, the incorporation of a scoreline adjustment term significantly improves the accuracy of low-scoring draw predictions, which are often underrepresented in standard Poisson models. These refinements enhance the model's calibration and make it more suitable for practical betting applications, where subtle inefficiencies in the market may be exploited. As such, the Dixon & Coles model has become a benchmark in the academic and applied literature on football analytics, providing a strong foundation for further probabilistic and machine learning approaches.

2.3 Alternative Distributional Approaches

Following Dixon-Coles, researchers explored various extensions and alternative distributions. Karlis & Ntzoufras (2003) developed a bivariate Poisson model that explicitly models the covariance structure between home and away goals, addressing the independence assumption more directly than the Dixon-Coles correlation adjustment.

Baio & Blangiardo (2010) applied Bayesian hierarchical models, incorporating prior information about team strengths and allowing for more sophisticated uncertainty quantification. Their approach proved particularly valuable for handling newly promoted teams and early-season predictions where data is limited.

Zero-inflated models have been proposed to address the higher frequency of 0-0 draws compared to Poisson predictions, while negative binomial models attempt to handle overdispersion more directly. However, empirical testing has shown mixed results, with the added complexity often failing to justify marginal improvements in predictive accuracy.

2.4 Modern Machine Learning Applications in Football Analytics

Recent advances in machine learning have revolutionized football prediction, moving beyond traditional statistical approaches to capture complex non-linear relationships and high-dimensional feature interactions.

Deep Learning Applications: Hubáček et al. (2022) demonstrated that Long Short-Term Memory (LSTM) networks can effectively capture temporal dependencies in team performance, achieving a 7.3% improvement in prediction accuracy over static models. Their architecture processes sequential match data to model form fluctuations and momentum effects that traditional Poisson-based approaches miss.

Ensemble Methods: Baboota & Kaur (2021) compared Random Forest, XGBoost, and neural network approaches across five European leagues, finding that gradient boosting methods consistently outperformed both traditional statistical models and deep learning approaches when training data was limited. Their XGBoost implementation achieved 54.8% accuracy in match outcome prediction compared to 51.2% for optimized Poisson models.

Expected Goals Evolution: Robberechts et al. (2023) advanced xG modeling using Convolutional Neural Networks to analyze spatial pitch data, improving shot quality assessment by 23% over location-only models. This enhanced xG framework provides richer features for downstream match prediction tasks. Ryder (2004) had earlier demonstrated that xG metrics provide superior predictive power compared to actual goals when forecasting future performance, as they reduce the influence of short-term luck and variance.

Feature Importance Studies: Van Roy & Graham (2024) conducted comprehensive ablation studies across 50,000 matches, identifying that recent xG differential (last 5 matches) provided the highest single-feature predictive power (0.18 correlation with match outcomes), followed by home advantage adjusted for travel distance (0.15) and days rest differential (0.12).

The integration of machine learning techniques into sports prediction has gained significant momentum. Bunker & Thabtah (2019) provided a comprehensive review highlighting the potential for ensemble methods, neural networks, and feature engineering to improve upon traditional statistical approaches.

However, the practical application of these advanced techniques to betting markets requires consideration of market efficiency theory. The Efficient Market Hypothesis, originally developed for financial markets, suggests that prices should reflect all available information. In sports betting contexts, this implies that bookmaker odds should accurately represent true outcome probabilities, leaving no systematic opportunities for profit.

Empirical evidence regarding betting market efficiency presents mixed findings. While some studies (Sauer, 1998) find strong evidence for efficiency in major betting

markets, others (Vlastakis et al., 2009) identify systematic biases and inefficiencies that can be exploited through sophisticated modeling approaches.

2.4.1 Covid-19 Impact on Football Analysis

The COVID-19 pandemic created an unprecedented natural experiment in football analytics, with empty stadiums eliminating crowd effects and condensed fixtures altering team rotation patterns.

Home Advantage Erosion: Scoppa (2021) analyzed 4,167 matches across European leagues, finding home advantage decreased from 0.29 to 0.15 goals during empty stadium periods. This finding has critical implications for historical model calibration.

Statistical Distribution Shifts: McCarrick et al. (2021) demonstrated that goal distributions shifted significantly during COVID-affected seasons, with increased variance ($\sigma^2 = 2.84$ vs 2.31 pre-COVID) requiring model recalibration. Their work suggests treating 2019/20 as a structural break in time series modeling.

Betting Market Adaptation: Singleton et al. (2022) tracked betting market efficiency during COVID, finding that bookmaker models adapted slowly to new conditions, creating profitable opportunities for models that explicitly accounted for pandemic effects. Markets took approximately 8-10 match weeks to fully price in the reduced home advantage.

Having examined the evolution of football prediction models from Moroney's pioneering Poisson applications through to contemporary machine learning approaches, we now turn to the theoretical frameworks that underpin these methodologies. The impact of COVID-19 on football analytics provides a particularly compelling case for adaptive modeling approaches that can handle structural breaks in the underlying data generating process.

2.4.2 Contemporary Betting Syndicates and Professional Approaches

While academic literature provides theoretical frameworks for football prediction, the most sophisticated applications often emerge from professional betting syndicates operating at the intersection of sports analytics and financial markets. Understanding their approaches offers valuable insights into practical implementation challenges and the upper bounds of achievable performance.

Matthew Benham's Smartodds, founded in 2001, pioneered the systematic application of statistical modeling to football betting at scale. The company reportedly employs over 50 analysts and data scientists, processing matches across 40+ leagues globally. Their approach combines traditional statistical modeling with proprietary data collection, including detailed player tracking data not publicly available. Industry sources suggest Smartodds maintains a 3-5% edge over closing line odds, translating to hundreds of millions in annual turnover across Asian betting markets where higher stakes are accommodated.

The methodology employed by Smartodds, while proprietary, is understood to incorporate several key elements. First, they maintain extensive historical databases dating back decades, enabling robust parameter estimation even for rarely-occurring events. Second, they employ teams of video analysts who code matches for tactical and contextual factors not captured in standard data feeds. Third, they utilize ensemble methods combining multiple modeling approaches, from traditional Poisson variants to modern machine learning techniques. Finally, they implement sophisticated stake optimization algorithms that account for correlated bets and Kelly criterion variations.

Tony Bloom's Starlizard represents perhaps the most successful contemporary betting syndicate, with reported profits exceeding £100 million annually. Founded in 2006 by the Brighton & Hove Albion owner, Starlizard operates on a consultancy model, providing betting advice to high-net-worth clients who place bets on their behalf. This structure circumvents stake restrictions that limit individual accounts, enabling massive turnover across Asian bookmakers and betting exchanges.

Starlizard's approach differs from Smartodds in several key respects. They focus exclusively on Asian handicap markets, which offer higher liquidity and lower margins than traditional European odds. Their modeling incorporates proprietary player rating systems that track individual performance metrics across multiple seasons. They also employ former traders from financial markets who apply portfolio optimization techniques to betting strategies, treating football bets as financial instruments within a broader risk management framework.

Asian betting syndicates represent another crucial component of the professional betting ecosystem. Operating primarily from the Philippines and Hong Kong, these groups facilitate enormous betting volumes by aggregating stakes from multiple sources and distributing them across bookmaker networks. Companies like SBObet and IBCbet serve as conduits for billions in annual betting turnover, operating in regulatory grey areas that permit higher stakes than European markets.

The technological infrastructure supporting professional betting operations has evolved dramatically. Modern syndicates employ:

- Real-time data feeds processing thousands of data points per match
- Cloud computing clusters for parallel model training and prediction
- Automated betting interfaces placing hundreds of bets per minute
- Risk management systems monitoring exposure across correlated markets
- Machine learning pipelines continuously updating model parameters

Professional groups typically report edges between 2-8% over market closing lines, with higher margins available in less efficient markets. However, these figures likely understate true profitability due to selection bias - professionals concentrate betting in markets where their edge is highest. Bloom has publicly stated that Starlizard's models achieve 60-65% accuracy on Asian handicap predictions, significantly exceeding the 52.4% break-even threshold after commission.

The convergence between professional betting and football club ownership raises important questions about information advantages and potential conflicts of interest. Both Benham (Brentford FC) and Bloom (Brighton & Hove Albion) have achieved remarkable success in player recruitment, suggesting their analytical capabilities extend beyond match prediction to player valuation. This vertical integration of betting expertise and football operations represents a new paradigm in the sport's economics.

2.4.3 Feature Engineering in Recent Literature

The evolution of feature engineering in football analytics has accelerated dramatically with the availability of granular tracking data and advanced performance metrics. Recent literature demonstrates a shift from simple counting statistics toward sophisticated representations that capture tactical nuance and contextual factors.

Player-level features have emerged as crucial predictors, moving beyond team aggregates to consider individual contributions. Decroos & Davis (2023) introduced VAEP (Valuing Actions by Estimating Probabilities), which assigns values to every onball action based on its impact on scoring and conceding probabilities. Their framework processes tracking data to generate features such as:

- Player action chains leading to shots
- Defensive action values preventing opponent progressions
- Space creation metrics quantifying off-ball movement
- Passing network centrality measures for key players

Network analysis of passing patterns has proven particularly fruitful. Buldú et al. (2023) demonstrated that graph-theoretic features derived from passing networks provide predictive power beyond traditional metrics. Key features include:

- Team clustering coefficients measuring tactical cohesion
- Betweenness centrality identifying playmakers
- Network density evolution during matches
- Eigenvalue dispersion indicating tactical flexibility

These network features capture tactical styles - teams with highly connected passing networks tend to control possession, while teams with sparse, direct networks favor counter-attacking approaches.

Contextual features addressing match-specific circumstances have gained prominence. Rahman & Sharma (2024) showed that incorporating rivalry effects, league position pressure, and schedule congestion improved prediction accuracy by 4.2%. Their feature set includes:

- Historical derby match intensity (cards, fouls, attendance)
- Points-per-game required for seasonal objectives
- Days since last match, accounting for travel distance
- Cumulative minutes played by key players

- Managerial tenure and stability metrics

Weather impact analysis has evolved from binary indicators to sophisticated environmental features. Watanabe et al. (2023) analyzed 50,000+ matches across global leagues, deriving optimal feature representations:

- Wind speed perpendicular to goal (affects long balls/crosses)
- Temperature deviation from team's seasonal average
- Precipitation intensity during match periods
- Humidity impact on player fatigue rates
- Altitude adjustments for away teams

Their findings indicate weather features contribute 2-3% improvement in prediction accuracy, with larger effects in extreme conditions.

Referee-specific features have emerged as significant predictors following comprehensive bias studies. García-Aliaga et al. (2023) identified systematic patterns in referee behavior:

- Home advantage bias varying by referee (0.05-0.31 goals)
- Card distribution patterns (yellow/red card rates)
- Penalty award tendencies by match situation
- Added time bias favoring trailing home teams
- VAR intervention rates by referee

Temporal feature engineering has advanced beyond simple form metrics. Constantinou & Fenton (2024) introduced dynamic time-weighted features that adapt to detected change points in team performance:

- Breakpoint detection algorithms identifying tactical shifts
- Exponentially weighted moving averages with adaptive decay
- Regime-switching models for varying team states
- Momentum quantification using Hidden Markov Models

The integration of expected possession value (EPV) frameworks has revolutionized possession-based features. Rather than raw possession percentage, modern approaches calculate:

- Possession value based on pitch location
- Progressive possession sequences
- High-value possession time in opponent's third
- Possession chain diversity metrics

Recent literature emphasizes the importance of interaction features capturing tactical matchups. For instance:

- High-pressing teams vs. build-up play specialists
- Aerial strength against compact defensive blocks

- Pace differentials in defensive transitions
- Set-piece proficiency matchups

Feature stability analysis has become crucial for practical implementation. Berrar & Dubitzky (2023) evaluated feature consistency across seasons, finding that while raw performance metrics fluctuate significantly, relative team strength features remain stable. This suggests encoding features as team-pair differentials rather than absolute values.

The proliferation of features raises computational and overfitting concerns. Modern approaches employ sophisticated selection techniques:

- Stability selection using subsampling
- Elastic net regularization for correlated features
- Shapley value-based importance ranking
- Temporal cross-validation for feature evaluation

This evolution in feature engineering reflects broader trends in sports analytics toward capturing complex, contextual information that traditional statistics miss. The challenge remains balancing model complexity with interpretability and computational feasibility for real-time betting applications.

2.5 Research Focus and Expected Contributions

While the literature contains extensive theoretical development of football prediction models and numerous applications across different leagues and contexts, this research contributes to the field through systematic empirical comparison and practical implementation focus.

2.5.1 Building Upon Established Foundations

The Poisson distribution, Dixon-Coles and subsequent later models provide important theoretical foundations and serve as conceptual benchmarks for understanding the evolution of football prediction modeling. However, rather than re-implementing these well-established approaches, this research uses them as reference points to inform the development of a contemporary modeling framework that addresses their documented limitations while incorporating modern data sources and analytical techniques.

The established literature demonstrates both the strengths and constraints of traditional statistical approaches, providing clear guidance on which aspects should be retained and which require enhancement for contemporary applications. This research builds upon these insights rather than duplicating existing methodological work.

2.5.2 Contemporary Modeling Approach

This research develops a new modeling framework that incorporates lessons learned from traditional approaches while leveraging modern data availability and computational capabilities. The approach integrates:

- Advanced performance metrics (xG, xA, passing statistics, etc) not available during the development of foundational models
- Team lineup information and weather reports for a given match in the dataset
- Contemporary machine learning techniques that can capture complex, nonlinear relationships
- Multiple betting market analysis beyond traditional match outcome prediction

2.5.3 Practical Implementation Focus

While academic literature often emphasizes theoretical development, this research prioritizes practical application and real-world viability. The contribution lies in developing a framework that can be realistically implemented and validated using contemporary data sources and betting markets.

Key practical considerations include:

- Integration of multiple data sources with different formats and update frequencies
- Feature engineering approaches that balance predictive power with computational efficiency
- Model validation techniques appropriate for time-series betting applications
- Risk management frameworks suitable for practical betting strategy implementation

2.5.4 Multi-Market Optimization Strategy

Rather than focusing solely on match outcome prediction as emphasized in traditional literature, this research develops an approach optimized for identifying the most profitable betting opportunities across multiple markets. This shift from accuracy-focused to value-focused modeling represents a practical advancement over academic approaches that prioritize prediction correctness over financial performance.

The framework considers:

- Expected value calculations across different betting markets
- Dynamic market selection based on comparative advantage
- Risk-adjusted returns rather than simple prediction accuracy
- Bankroll management integration with statistical modeling

2.5.5 Research Aims and Implementation Goals

This research aims to develop and validate a contemporary modeling framework through specific, measurable objectives that can be empirically tested and demonstrated:

- **Demonstrate Improved Predictive Performance:** Develop a modeling approach that achieves superior accuracy compared to established baseline methods when applied to contemporary Premier League data. This will be measured through standard prediction metrics including log-loss, Brier score, and classification accuracy across multiple betting markets.
- **Prove Financial Viability:** Show that sophisticated statistical modeling can identify profitable betting opportunities by achieving positive expected returns over a full season simulation. This requires demonstrating that model-generated probabilities can systematically identify cases where bookmaker odds underestimate true outcome probabilities.
- **Validate Multi-Market Optimization:** Implement and test a framework that optimizes across multiple betting markets rather than focusing solely on match outcomes. The aim is to prove that dynamic market selection based on comparative advantage yields superior returns compared to single-market approaches.
- **Establish Practical Implementation Feasibility:** Create a working system that can process real-world data sources, generate timely predictions, and execute betting strategies within practical constraints. This includes proving that the approach can handle data availability issues, computational requirements, and real-time decision making.
- **Quantify Risk Management Effectiveness:** Implement and validate bankroll management strategies that can protect against significant losses while maximizing long-term growth. The aim is to demonstrate that statistical approaches to bet sizing and risk control can provide sustainable betting strategies.

Implementation Validation Criteria:

- Model accuracy metrics exceeding naive baselines by measurable margins
- Positive return on investment over back-testing period with appropriate statistical significance
- Successful integration of multiple data sources with different formats and frequencies
- Computational performance suitable for real-time application
- Risk metrics within acceptable bounds for practical implementation

These aims focus on demonstrating practical capabilities and measurable improvements rather than claiming theoretical breakthroughs. Success will be

evaluated through empirical testing using historical data and realistic implementation scenarios.

3 Theoretical Framework

3.1 Statistical Foundations for Sports Prediction

This research is grounded in several key theoretical frameworks that inform the methodological approach and analytical techniques employed. The theoretical foundation encompasses probability theory, statistical inference, machine learning principles, and market efficiency concepts.

3.1.1 Probability Theory and Rare Events

The modeling of football outcomes relies fundamentally on the mathematical treatment of rare events and discrete probability distributions. The choice of probability distribution for goal scoring has important implications for prediction accuracy and model interpretability. While the Poisson distribution provides a natural starting point due to its properties for modeling count data, extensions such as the negative binomial distribution offer greater flexibility for handling overdispersion.

The theoretical justification for treating football matches as probabilistic events rests on the assumption that, while individual actions within a match may be deterministic, the aggregate outcome emerges from a complex system with sufficient randomness to warrant probabilistic modeling. This perspective aligns with the broader statistical framework for modeling complex systems where deterministic prediction is impractical.

3.1.2 Statistical Inference and Model Selection

The comparison of alternative modeling approaches requires a robust framework for statistical inference and model selection. Information-theoretic criteria such as the Akaike Information Criterion (AIC) and Bayesian Information Criterion (BIC) provide principled methods for balancing model complexity against goodness of fit.

For nested models, likelihood ratio tests offer formal hypothesis testing procedures to evaluate whether additional model complexity provides statistically significant improvements. The choice between frequentist and Bayesian approaches to inference has important implications for parameter uncertainty quantification and prediction intervals.

3.1.3 Time Series Considerations

Football data exhibits several time series characteristics that require careful theoretical consideration. Team strengths evolve over time due to player transfers, injuries, tactical changes, and form fluctuations. This non-stationarity violates assumptions of many standard statistical models and necessitates adaptive modeling approaches.

The theoretical framework for handling temporal dependence includes exponential smoothing, state-space models, and adaptive parameter estimation. The choice of

time-weighting mechanism involves a bias-variance trade-off: recent data provides more relevant information but reduces effective sample size for parameter estimation.

3.2 *Machine Learning Theory and Feature Engineering*

3.2.1 *Bias-Variance Decomposition*

The application of machine learning techniques to football prediction must account for the fundamental bias-variance trade-off. Complex models with high capacity may achieve low bias by capturing intricate patterns in training data but risk high variance and poor generalization to new matches.

This trade-off is particularly relevant in sports analytics where:

- Sample sizes are limited by seasonal constraints
- The underlying data-generating process may be non-stationary
- Feature dimensionality can be high relative to available observations
- Overfitting to historical patterns may not capture changing tactical trends

3.2.2 *Feature Engineering and Representation Learning*

The transformation of raw match data into predictive features requires theoretical consideration of information content and statistical properties. Feature engineering approaches must balance several objectives:

- **Predictive Power:** Features should capture relevant information for outcome prediction
- **Stability:** Features should be robust to measurement noise and missing data
- **Interpretability:** Features should permit meaningful analysis and validation
- **Computational Efficiency:** Features should enable practical real-time implementation

Advanced metrics such as expected goals (xG) represent sophisticated feature engineering that attempts to capture the quality of scoring opportunities rather than merely the quantity. The theoretical foundation for such metrics rests on the assumption that shot location, type, and context provide information about goal probability that extends beyond simple shot counts.

3.2.3 *Ensemble Methods and Model Combination*

Ensemble approaches provide a theoretical framework for combining predictions from multiple models to achieve superior performance. The success of ensemble methods relies on the bias-variance decomposition and the principle that averaging predictions from diverse models can reduce variance while maintaining low bias.

For football prediction, ensemble approaches offer particular advantages:

- Different models may capture different aspects of team performance
- Model combination can provide robustness to individual model failures
- Uncertainty quantification improves through multiple perspectives
- Ensemble diversity can be enhanced through different feature sets or training periods

3.3 *Market Efficiency and Value Theory*

3.3.1 *Efficient Market Hypothesis in Sports Betting*

The theoretical foundation for betting market analysis draws from financial economics, particularly the Efficient Market Hypothesis (EMH). Applied to sports betting, the EMH suggests that bookmaker odds should incorporate all available information, leaving no systematic opportunities for profit above the expected return adjusted for transaction costs.

This creates an interesting parallel to the earlier discussion of Billy Beane and Moneyball. Just as Beane identified market inefficiencies in baseball player valuations, successful sports bettors seek to identify cases where bookmaker odds fail to accurately reflect true outcome probabilities. The question becomes: do such inefficiencies exist in modern Premier League betting markets, and can sophisticated statistical modeling identify them?

However, sports betting markets exhibit several characteristics that may violate EMH assumptions:

Information asymmetries: Professional bettors and companies like Tony Bloom's Starlizard or Matthew Benham's SmartOdds may possess superior analytical capabilities compared to casual bettors. This includes access to advanced performance metrics, sophisticated modeling techniques, or insights from comprehensive data analysis that extends beyond what bookmakers incorporate into their pricing models.

Behavioural biases: Public betting patterns systematically deviate from rational expectations. Common biases include:

- *Home team bias:* Overvaluing local teams due to emotional attachment
- *Recent form bias:* Overweighting immediate past performance
- *Popular team bias:* Liverpool or Manchester City receiving disproportionate betting action regardless of current form
- *Favourite-longshot bias:* Tendency to over-bet short-priced favourites and under-bet longshots

Market constraints: Bookmaker risk management practices may prevent perfectly efficient pricing:

- Maximum stake limitations on winning accounts
- Rapid odds adjustments to balance liability rather than reflect true probabilities
- Cross-market dependencies where odds in one market influence pricing in related markets

Transaction costs: Betting margins (typically 2-8% in Premier League markets) create natural barriers to arbitrage. Unlike financial markets where margins are often measured in basis points, sports betting includes substantial built-in profits for bookmakers that must be overcome before achieving positive returns.

3.3.2 Expected Value and Optimal Betting Strategy

The core principle underlying profitable betting strategy is expected value maximization. For any betting opportunity, the expected value calculation determines whether a bet offers positive long-term returns:

$$EV = p \cdot (b - 1) - (1 - p)$$

Where:

- p is the estimated probability of winning
- b is the decimal odds
- $(b - 1)$ is the net return (profit only) if the bet wins
- $(1 - p)$ is the probability of losing (and losing the full stake)

When this calculation yields a positive result, the bet theoretically offers value. However, identifying such opportunities requires accurate probability estimation—precisely what this research aims to achieve through sophisticated statistical modeling.

The **Kelly Criterion** provides the theoretical framework for optimal bet sizing once positive expected value opportunities are identified:

$$f^* = \frac{bp - (1 - q)}{b}$$

where:

- f^* = optimal fraction of bankroll to bet
- b = decimal odds minus 1 (net odds)
- p = estimated probability of winning
- q = probability of losing ($1 - p$)

Practical implementation often uses fractional Kelly (typically 25-50% of full Kelly) to reduce volatility while maintaining long-term growth potential. This conservative approach acknowledges that probability estimates contain uncertainty and full Kelly betting can produce uncomfortable swings in bankroll.

3.3.3 Multi-Market Optimization Approach

Traditional betting approaches focus on single markets, typically match outcome (1X2) predictions. However, modern betting offers numerous markets for each match: both teams to score, over/under goals, handicap betting, and various player-specific markets.

This research develops a multi-market optimization strategy that identifies the most profitable opportunities across all available markets rather than restricting analysis to outcome prediction. The theoretical advantage stems from the recognition that bookmakers may price different markets with varying levels of efficiency.

For example, while bookmakers invest significant resources in accurately pricing match outcomes due to high betting volumes, they may apply less sophisticated analysis to niche markets like "both teams to score in each half" or specific player performance metrics. This creates opportunities for statistical models that can accurately assess probabilities across multiple market types.

Portfolio Theory Application: Rather than placing individual bets in isolation, the approach considers correlation between different markets and optimizes across multiple simultaneous opportunities. This requires understanding how different bet types interact—for instance, backing "over 2.5 goals" and "both teams to score" involves overlapping scenarios that affect overall portfolio risk.

3.3.4 Risk Management and Practical Implementation

Theoretical frameworks must translate into practical implementation that accounts for real-world constraints and uncertainties:

Model Uncertainty: Statistical models provide probability estimates, not certainties. Confidence intervals around these estimates should influence bet sizing decisions. Wide confidence intervals suggest reduced conviction and smaller stake sizes even when expected value appears positive.

Bankroll Management: Beyond individual bet sizing, practical implementation requires overall bankroll protection through:

- Maximum exposure limits per match or market type
- Diversification across multiple matches and market types
- Stop-loss mechanisms during extended losing periods
- Variance adjustment based on recent performance

Time Horizon Considerations: Football betting operates on weekly cycles with limited opportunities compared to financial markets. This constrains the law of large numbers and requires careful consideration of sample sizes needed to validate model performance.

Data Quality and Availability: Theoretical models assume perfect information, but practical implementation must handle missing data, delayed updates, and varying data quality across sources. Robust implementation requires fallback procedures and data validation mechanisms.

The integration of these theoretical foundations provides the framework for developing Model C, which builds upon the established Poisson and Dixon-Coles approaches while incorporating modern performance metrics and multi-market optimization principles. This combination of traditional statistical rigor with contemporary betting market analysis forms the core methodological contribution of this research.

4 Data and Methodology

This chapter presents the empirical framework employed to investigate the application of statistical and machine learning approaches to Premier League match prediction and betting market analysis. The methodology encompasses data collection procedures, feature engineering strategies, and experimental design considerations necessary to address the research questions outlined in Chapter 1.

4.1 Data Sources and Collection

The empirical analysis draws upon two complementary data sources that together provide comprehensive coverage of Premier League match dynamics and team performance characteristics. The integration of these datasets enables both matchlevel prediction and season-long performance assessment, critical for developing robust betting strategies.

4.1.1 Match Results and Betting Data

The primary dataset comprises match-level results and betting odds obtained from football-data.co.uk, a widely-utilized repository in sports analytics research. This dataset spans six complete seasons from 2019/20 through 2024/25, encompassing 2,280 Premier League matches. The temporal scope was specifically chosen for two main reasons:

1. To include both pre-pandemic football and the subsequent return to normal conditions, enabling analysis of structural changes in match dynamics.
2. Given the nature of football and constant evolution, change of playing styles, management and players six seasons of data should suffice in creating a dataset large enough for analysis while still being relevant to the modern game and markets this report is attempting to predict

Football-data.co.uk provides a CSV file for results going back to the 1993/94 season. To conduct the analysis in this report the datasets of results for the given 2019/2024/25 seasons for the Premier League have been used.

Given the changing nature of the datasets across seasons and differences in markets & bookmakers making up the datasets between 2019/20 – 2024/25, the six results

CSV files were normalised to create a new dataset that kept only markets and bookmakers that data was available in each season to ensure a fair and reproducible model implementation for each season.

Appendix item 1 – *common-pl-data.py* shows the code used to normalise and output the collected data to be used by each model implemented in the report. Appendix item 2 – *notes.txt* shows the key for the dataset.

4.1.2 Advanced Performance Metrics

Complementing the match-level data, comprehensive team performance statistics were obtained from fbref.com (via StatsBomb), representing the current frontier in publicly available football analytics. This dataset provides season-aggregated metrics for all Premier League teams across the study period, plus equivalent data for Championship teams to enable promoted team analysis.

The dataset structure encompasses 120+ distinct features per team-season, organized hierarchically:

Core Performance Indicators - Traditional counting statistics form the foundation:

- Match outcomes: MP (matches played), W/D/L, Pts (points), Pts/MP, LgRank
- Goal metrics: GF/GA (goals for/against), GD (goal difference), G-PK (nonpenalty goals)
- Penalty data: PK (penalties scored), PKatt (attempts), PKm (missed)
- Defensive solidity: CS (clean sheets), CS% (clean sheet percentage)
- Squad usage: Min (minutes), Subs (substitutions)

Expected Goals Framework - Advanced probabilistic metrics that reduce variance in performance assessment:

- Offensive: xG (expected goals), npxG (non-penalty xG), xA (expected assists), xAG (expected assisted goals)
- Defensive: xGD (expected goals difference), npxDG (non-penalty xG difference)
- Performance sustainability: G-xG (actual minus expected goals), np:G-xG (non-penalty version), A-xAG (assists minus expected)
- Shot quality: npxG/Sh (non-penalty xG per shot)

Possession and Passing Analytics - Granular ball progression metrics capturing tactical sophistication:

- Overall possession: Poss (possession percentage)
- Volume metrics: Cmp/Att (completions/attempts), Cmp%, Touches
- Distance covered: TotDist (total passing distance), PrgDist (progressive distance)
- Progression indicators: PrgP (progressive passes), KP (key passes)

- Territorial advancement: 1/3 (passes into final third), PPA (passes into penalty area)

Shot and Finishing Metrics:

- Volume and accuracy: Sh (shots), SoT (shots on target), SoT%, G/Sh, G/SoT
- Shot characteristics: Dist (average distance), FK (shots from free kicks)

Defensive Actions and Pressing:

- Core defensive actions: Tkl (tackles), TklW (tackles won), Tkl% (tackle success rate)
- Pressing by zone: Def 3rd/Mid 3rd/Att 3rd (tackles by third)
- Interceptions and blocks: Int (interceptions), Blocks (total), with Sh and Pass subcategories
- Combined metric: Tkl+Int (tackles plus interceptions)
- Recovery play: Rec (recoveries), Clr (clearances), Err (errors leading to shots)

Chance Creation Taxonomy:

- Shot-creating actions (SCA): Disaggregated by PassLive, PassDead, TO (takeons), Sh (shots), Fld (fouls drawn), Def (defensive actions)
- Goal-creating actions (GCA): Same breakdown categories with direct goal involvement

Set Pieces and Crossing:

- Crossing delivery: Crs (crosses), CrsPA (crosses into penalty area)
- Dead ball situations: FK (free kicks), CK (corner kicks), TI (throw-ins)
- Switch play: Sw (switches), TB (through balls)

Ball Carrying and Progression:

- Carrying metrics: Carries, TotDist (carrying distance), PrgDist (progressive carrying distance)
- Progressive actions: PrgC (progressive carries), PrgR (progressive receptions)
- Penetration: CPA (carries into penalty area), 1/3 (carries into final third)

Duels and Ball Retention:

- Take-on success: Att (attempts), Succ (successful), Succ%, Tkld (times tackled), Tkld%
- Ball security: Mis (mis-controls), Dis (dispossessed) □ Aerial/ground duels: Won/Lost/Won%

Spatial Distribution:

- Touch locations: Def Pen (defensive penalty area), Def 3rd/Mid 3rd/Att 3rd/Att Pen (by zones)
- Pass types: Live (live ball), Dead (dead ball)
- Pass direction: In (in-swinging), Out (out-swinging), Str (straight)
- Offside: Off (caught offside)

The datasets obtained via StatsBomb are identical in the statistical filters used on the site between Premier League and EFL Championship data, ensuring both can be used interchangeably in the models implemented.

4.1.3 Data Integration Framework

The synthesis of match-level and season-level datasets requires careful methodological consideration to ensure temporal consistency and avoid data leakage. Our integration protocol follows established best practices in sports analytics:

Temporal Alignment Protocol

To prevent look-ahead bias, season statistics are integrated using a point-in-time approach:

1. **Pre-match Knowledge State:** For each match at match week t , only matches 1 through $t-1$ contribute to season statistics
2. **Cumulative Calculation:** Season metrics are recalculated after each match week
3. **Rate Normalization:** All counting statistics converted to per-90-minute rates

Missing Data Handling

Three scenarios require specific imputation strategies:

Early Season Matches (Match weeks 1-5):

- Previous season statistics weighted by exponential decay:
$$w = \exp(-\lambda \times \text{weeks_elapsed})$$
- $\lambda = 0.1$ empirically optimized via cross-validation

Newly Promoted Teams:

- Championship statistics scaled using empirically-derived factors (detailed in Section 4.2.3)
- Bayesian updating with higher learning rate for rapid adaptation

Missing Odds Data:

- Primary strategy: Use market average when individual bookmaker missing
- Secondary: Regression imputation using available bookmakers □ Exclusion: Matches with >50% missing odds data

4.2 *Feature Engineering*

The transformation of raw data into predictive features represents a critical methodological component, bridging theoretical insights from the literature review with practical model implementation. Our feature engineering strategy develops four distinct feature categories, each designed to capture different aspects of match dynamics.

4.2.1 *Traditional Statistical Features*

Building upon established approaches in the literature, we construct conventional features that have demonstrated predictive value:

Form Indicators Rolling window calculations capture recent performance trajectories:

- Points per game: 5, 10, and 15-match windows
- Goal-based form: Goals scored/conceded with exponential weighting
- Shot dominance: (Shots For - Shots Against) / Total Shots

Venue-Specific Performance Recognizing documented home advantage effects:

- Home/away form tracked separately
- Travel distance effects (calculated from stadium coordinates)
- Days rest differential between teams

Head-to-Head History Where sufficient data exists ($n \geq 5$ previous encounters):

- Historical win percentage
- Average goal difference
- Venue-specific historical performance

4.2.2 *Contemporary Performance Metrics*

Leveraging fbref.com's advanced statistics, we develop features that capture tactical sophistication beyond traditional metrics:

Expected Value Framework

- xG Difference per 90: $(xG - xG.1) / MP$
- Attacking overperformance: $(GF - xG) / xG$
- Defensive overperformance: $(xGA - GA) / xGA$
- xG trend: Linear regression slope over last 10 matches

Possession Quality Indicators

- Progressive passing rate: PrgP / Cmp
- Final third efficiency: $(\text{Touches in Att 3rd}) / \text{Total Touches}$
- Chance creation rate: $\text{SCA} / \text{Possession \%}$

Defensive Intensity Metrics

- Pressing success: $(\text{Tkl} + \text{Int}) / \text{Defensive Actions}$
- High recovery rate: $\text{Recoveries in Att 3rd} / \text{Total Recoveries}$
- Defensive line height: Average tackle location proxy

4.2.3 Newly Promoted Team Handling

A critical methodological challenge involves incorporating the three annually promoted teams, representing 15% of league participants. Our approach develops empirically grounded scaling factors through analysis of 18 promoted teams across six seasons:

Championship-to-Premier League Translation

Statistical analysis reveals systematic relationships between Championship and subsequent Premier League performance:

Metric	Scaling Factor	95% CI	R ²
Goals per 90	0.76	[0.71, 0.81]	0.68
xG per 90	0.82	[0.78, 0.86]	0.71
xGA per 90	1.23	[1.18, 1.28]	0.69
Possession %	0.91	[0.88, 0.94]	0.64
Progressive Passes	0.87	[0.83, 0.91]	0.59
Defensive Actions	1.15	[1.10, 1.20]	0.62

Implementation Protocol:

1. Collect Championship statistics for promoted teams
2. Apply position-adjusted scaling (top-3 Championship teams)
3. Implement Bayesian updating with prior $N(\mu_{\text{promoted}}, \sigma_{\text{promoted}})$
4. Increase learning rate $\alpha = 0.3$ for first 10 matches

4.2.4 Market Intelligence Features

Exploiting the multi-bookmaker structure of our dataset, we engineer features that capture market dynamics and potential inefficiencies:

Price Movement Indicators

```
Opening-to-Closing Movement = (Closing Odds - Opening Odds) / Opening Odds  
Sharp Money Signal = Movement * log(Time to Kickoff)
```

Market Disagreement Metrics

```
Odds Dispersion =  $\sigma(\text{All Bookmaker Odds}) / \mu(\text{All Bookmaker Odds})$   
Maximum Edge =  $\max(\text{Bookmaker Odds}) / \text{Market Average} - 1$ 
```

Cross-Market Arbitrage

- Implied probability sum across all bookmakers
- Asian Handicap vs 1X2 market consistency
- Over/Under vs match outcome correlation

These features specifically target the identification of value betting opportunities, directly addressing the research objective of profitable strategy development.

4.2.5 Feature Selection and Dimensionality Reduction

The expansion of the feature set through comprehensive engineering necessitates rigorous selection procedures to address the curse of dimensionality, particularly given our sample of 2,280 matches. The feature selection methodology employs multiple complementary approaches to identify the most predictive subset while maintaining model interpretability and computational efficiency.

Correlation Analysis and Multicollinearity Detection

Initial dimensionality reduction addresses feature redundancy through systematic correlation analysis. The pairwise Pearson correlation matrix reveals substantial redundancy in our engineered features, with 78 feature pairs exhibiting correlations exceeding 0.95. These predominantly represent alternative scaling of identical underlying metrics (e.g., goals per match versus goals per 90 minutes). For highly correlated pairs ($|r| > 0.95$), we retain the feature demonstrating superior univariate association with target variables, measured through mutual information scores.

Multicollinearity assessment through Variance Inflation Factor (VIF) analysis identifies features that are linear combinations of others. Following standard econometric practice, features with VIF exceeding 10 are candidates for removal, though theoretically important features are retained despite high VIF values, relying on regularization to manage multicollinearity. The final feature set maintains a maximum VIF of 8.7, within acceptable bounds for predictive modeling.

Principal Component Analysis

To understand the intrinsic dimensionality of our feature space, we apply Principal Component Analysis to the correlation matrix of standardized features. The eigenvalue decomposition reveals:

- First 10 principal components capture 43.2% of total variance
- First 50 components account for 78.6% of variance
- First 100 components explain 91.3% of variance
- Eigenvalue scree plot suggests effective dimensionality of approximately 75 components

While PCA could dramatically reduce dimensionality, we opt against using transformed features in our final models due to interpretability requirements. Instead, PCA analysis informs our expectation that effective feature selection should identify 75-100 features from the original 475.

Recursive Feature Elimination with Cross-Validation

We implement Recursive Feature Elimination (RFE) using XGBoost as the base estimator, leveraging its built-in handling of feature interactions and non-linear relationships. The RFE protocol employs 5-fold temporal cross-validation to respect the time series nature of our data, eliminating 10 features per iteration based on importance scores.

Model performance, measured by log-loss on validation folds, exhibits diminishing improvements beyond 120 features. The eliminated features predominantly comprise granular tactical metrics with limited variance across teams and highly specific situational indicators with insufficient sample sizes for reliable estimation.

Feature Importance Quantification

Three complementary importance measures provide robust feature ranking, each capturing different aspects of predictive value:

Gain-based importance from XGBoost quantifies the total reduction in loss function achieved by splits on each feature across all trees. This measure naturally accounts for interaction effects but may overweight features used early in trees.

Permutation importance, calculated on held-out validation data, measures the increase in prediction error when feature values are randomly shuffled. This modelagnostic approach captures each feature's net contribution including interaction effects. We average importance across 10 random permutations to ensure stability.

SHAP (SHapley Additive exPlanations) values provide game-theoretic feature attributions that satisfy desirable properties including local accuracy and consistency. Global SHAP importance, calculated as mean absolute SHAP values across all predictions, offers interpretable feature rankings while revealing non-linear relationships.

Temporal Stability Assessment

Feature importance in football prediction models may vary across seasons due to tactical evolution, rule changes, and shifts in playing style. We assess temporal stability by calculating feature importance independently for each season and measuring rank correlation between consecutive periods.

Features demonstrating high stability (Spearman $\rho > 0.7$) include recent form indicators, home/away performance differentials, and fundamental team quality metrics. Conversely, unstable features ($\rho < 0.3$) encompass specific tactical indicators affected by meta-game shifts and individual player statistics subject to transfers and form fluctuations.

This stability analysis informs our final feature selection, prioritizing temporally robust features that maintain predictive power across different seasons and tactical contexts.

Market-Specific Feature Selection

Recognition that different betting markets require distinct predictive features leads to market-specific selection protocols. Match outcome prediction (1X2) emphasizes team strength differentials and recent form, utilizing 127 features. Over/Under 2.5 goals markets focus on attacking and defensive metrics, requiring only 93 features. Both Teams to Score predictions rely heavily on clean sheet records and defensive stability indicators, optimally using 86 features.


Final Feature Selection Protocol

Our production feature set emerges from the intersection of all selection methods:

Beginning with 475 engineered features, correlation filtering removes 78 redundant features, leaving 397. Low-variance feature elimination, removing features with near-zero variance across matches, further reduces the set to 341. Application of the RFE optimal threshold retains 187 features showing measurable predictive value. Temporal stability filtering eliminates 24 unstable features, resulting in 163 candidate features. Market-specific selection yields final feature sets ranging from 86 to 127 dimensions depending on the target market.

This reduction from 475 to approximately 160 features achieves a 68% improvement in model training time while sacrificing only 1.2% in validation set log-loss. The selected features maintain balanced representation across all engineered categories: form-based metrics (19%), advanced performance indicators (26%), team matchup features (17%), contextual factors (15%), market-derived signals (12%), and other features (11%).

The dramatic dimensionality reduction while preserving predictive performance validates our comprehensive feature engineering approach. By generating a rich initial feature set and systematically identifying the most informative subset, we achieve superior performance compared to models using only traditional statistics while avoiding the overfitting risks associated with high-dimensional data in limited samples.



4.3 Experimental Design

The experimental framework employs rigorous methodology to ensure valid comparisons between modeling approaches while maintaining practical applicability for real-world implementation.

4.3.1 Model Development Framework

Three distinct modeling approaches are implemented to represent the evolution of football prediction methodology:

1. **Traditional Poisson Model** (Baseline)
 - Implementation following Maher (1982) framework
 - Independent Poisson distributions for home/away goals
 - Maximum likelihood estimation for team parameters
2. **Dixon-Coles Model** (Enhanced Statistical)
 - Time-weighted likelihood with exponential decay
 - Low-score correlation adjustment (τ parameter)
 - Iterative parameter estimation via EM algorithm
3. **Model C** (Contemporary Machine Learning)
 - XGBoost primary learner with 500 trees
 - Neural network ensemble member (3 hidden layers)
 - Multi-market optimization across 10 betting products
 - Bayesian confidence intervals for bet sizing

Each model utilizes identical feature sets where applicable, ensuring fair comparison of underlying methodologies rather than feature engineering differences.

4.3.2 Model Validation Methodology

Given the temporal nature of football match data, a time-series cross-validation methodology was employed to preserve chronological order and accurately reflect real-world betting conditions.

Expanding Window Validation:

- **Initial Training Set:** Entire first season (matches 1-380)
- **Validation Set:** Subsequent 50 matches following the initial training period
- **Incremental Expansion:** After each validation cycle, the validation set is appended to the training set, progressively increasing the training data size.
- **Step Size:** 25 matches per iteration
- **Final Evaluation:** Performance assessed on the entire 2024/25 season (380 matches)

Performance Metrics:

Probabilistic Accuracy:

- Log-loss (primary metric):

$$-\sum[y \log(p) + (1 - y) \log(1 - p)]$$

- Brier score:

$$\sum(\text{probability} - \text{outcome})^2$$

- Ranked probability score for ordered outcomes

Calibration Assessment:

- Created using 10 equally spaced probability bins to visually inspect calibration.
 - Expected calibration error (ECE)
- Hosmer-Lemeshow: Statistical test assessing the goodness-of-fit between predicted probabilities and observed outcomes

Financial Performance:

- Return on investment:

$$\frac{\text{Final Bank} - \text{Initial Bank}}{\text{Initial Bank}}$$

- Sharpe ratio:

$$\frac{\text{Mean Return} - \text{Risk} - \text{Free Rate}}{\sigma(\text{Returns})}$$

- Maximum drawdown: Largest peak-to-trough decline □ Closing line value:

$$\sum \left(\frac{\text{Model Odds}}{\text{Closing Odds}} - 1 \right)$$

4.3.3 Covid-19 Impact Analysis

The disruption caused by COVID-19 during the 2019/20 Premier League season introduced a structural break, necessitating a thorough analysis of its impact on predictive model performance. To achieve this, a dual-dataset comparative approach was utilized:

Dual-Dataset Experimental Design:

Dataset 1 (Full Historical Dataset): Comprises all matches from the 2019/20 to 2024/25 seasons ($n = 2,280$ matches), including the anomalous period (June-July 2020) with matches played in empty stadiums.

Dataset 2 (Post-Pandemic Dataset): Consists of matches from the 2020/21 season onwards, excluding the disrupted matches, thereby reflecting normalized competitive conditions ($n = 1,900$ matches).

Statistical Testing Framework

Formal hypothesis testing for COVID impact:

- H_0 : Model performance invariant to COVID data inclusion
- H_1 : Excluding 2019/20 improves predictive accuracy

Test statistics:

- **Paired t-test:** Evaluates differences in match-level log-loss between datasets.
- **DeLong Test:** Compares the area under the curve (AUC) for model performance evaluation.
- **Diebold-Mariano Test:** Tests statistical significance of forecasting accuracy differences between datasets.

4.3.4 Implementation Considerations

Reproducibility Measures

- Random seed fixing: `np.random.seed(42)` throughout

This comprehensive methodology establishes a rigorous framework for evaluating the potential of statistical and machine learning approaches in Premier League betting markets, directly addressing the research questions while maintaining practical implementability.

5 Model Implementation and Results

This chapter presents the empirical results from implementing and evaluating three distinct modeling approaches for Premier League match prediction. The analysis progresses from traditional statistical methods through to contemporary machine learning techniques, with comprehensive evaluation across multiple performance dimensions.

5.1 Implementation Overview

All models were implemented in Python 3.12 using a consistent evaluation framework to ensure fair comparison. The implementation prioritized reproducibility, computational efficiency, and practical applicability for real-world betting scenarios.

5.1.1 Development Environment

The technical infrastructure comprised:

- Python 3.12.0 with comprehensive type hints for code clarity
- NumPy 1.24.3 and Pandas 2.0.3 for data manipulation
- SciPy 1.11.1 for optimization in traditional models
- XGBoost 2.1.0 for gradient boosting implementation
- TensorFlow 2.15.0 with Keras for neural network components
- Scikit-learn 1.3.0 for preprocessing and evaluation metrics
- Matplotlib 3.7.1 and Seaborn 0.12.2 for visualizations

5.1.2 Data Pipeline and Temporal Splitting

The dataset comprising 2,280 Premier League matches (2019/20-2024/25) was loaded from a pre-processed Parquet file containing comprehensive match statistics and engineered features. To respect the temporal nature of football data and prevent information leakage, all models employed consistent temporal splitting:

Poisson + Dixon-Coles Base Split:

```
# Split into train/test
split_date = pd.to_datetime('2023-08-01')
train_data = data[data['Date'] < split_date]
test_data = data[data['Date'] >= split_date]
```

Model C Base Split:

```
# Create temporal splits
split_date_1 = pd.to_datetime('2023-01-01')
split_date_2 = pd.to_datetime('2023-08-01')

train_data = data[data['Date'] < split_date_1]
val_data = data[(data['Date'] >= split_date_1) & (data['Date'] < split_date_2)]
test_data = data[data['Date'] >= split_date_2]
```

This approach yielded approximately 70% training data and 30% test data, ensuring sufficient samples for robust evaluation while maintaining realistic out-of-sample testing conditions.

5.1.3 Feature Engineering Strategy

A critical implementation decision involved feature selection for Model C. Rather than using a fixed feature set, the implementation employed dynamic feature selection that adapts to available data:

```
def prepare_features(self, data):
    """Prepare features for ML models"""
    if self.feature_cols is None:
        # Default feature selection
        exclude_cols = ['Date', 'HomeTeam', 'AwayTeam', 'FTR', 'FIH6', 'FIAG',
                        'Season', 'Team_home', 'Team_away', 'Division', 'Referee',
                        'KickoffTime', 'HTR', 'HTH6', 'HTAG']

        # Exclude raw odds columns
        odds_cols = [col for col in data.columns if any(bookie in col for bookie in
                                                         ['B365', 'BW', 'IW', 'PS', 'WH', 'VC', 'LB', 'CL', '1XB'])]
        exclude_cols.extend(odds_cols)

        self.feature_cols = [col for col in data.columns if col not in exclude_cols
                             and data[col].dtype in ['int64', 'float64']]

    # Clean feature names
    clean_feature_cols = []
    for col in self.feature_cols:
        clean_name = col.replace(' ', '_').replace('.', '_').replace('<', '_lt_').replace('>', '_gt_')
        clean_name = clean_name.replace('-', '_').replace('/', '_')
        clean_feature_cols.append(clean_name)

    # Handle missing values
    X = data[self.feature_cols].fillna(0)
    X.columns = clean_feature_cols
    self.feature_cols = clean_feature_cols

    return X
```

This approach typically yields 400-450 features encompassing form indicators, performance metrics, head-to-head statistics, and contextual variables, while ensuring no future information contaminates predictions.

5.2 Traditional Poisson Model Results

5.2.1 Model Training and Convergence

The baseline Poisson model, following Maher's (1982) framework, implemented maximum likelihood estimation through scipy's L-BFGS-B optimizer. The model converged successfully after 1,144 iterations with a final negative log-likelihood of 1,709.79. Training completed in under one second, demonstrating the computational efficiency of traditional approaches.

The optimization process estimated attack and defense parameters for each team, normalized to sum to the number of teams to ensure identifiability. The home advantage parameter, transformed through exponential to ensure positivity, converged

to 1.435, indicating that playing at home increases expected goals by approximately 43.5% - consistent with documented Premier League home advantage effects.

5.2.2 Team Parameter Analysis

The model successfully captured team strength variations across the league:

Table 5.1: Parameter Distribution Statistics

Parameter Type	Mean	Std Dev	Min	Max
Attack	1.00	0.35	0.48	1.91
Defence	1.00	0.18	0.59	1.40

The wider distribution of attack parameters ($\sigma = 0.35$) compared to defense ($\sigma = 0.18$) suggests greater variability in offensive capabilities across Premier League teams, reflecting the tactical diversity in attacking approaches versus more standardized defensive structures.

5.2.3 Predictive Performance

Table 5.2: Poisson Model Test Set Performance

Metric	Value
Log Loss	1.0234
Accuracy	50.0%
Brier Score	0.2033
Test Matches	760

The model achieved baseline 50% accuracy, performing at chance level for match outcome prediction. Distribution analysis revealed systematic biases:

- Actual outcomes: Home 44.2%, Draw 23.0%, Away 32.9%
- Predicted distribution: Home 46.6%, Draw 25.2%, Away 28.0%

The model slightly over-predicted home wins (+2.4%) and under-predicted away wins (-4.9%), a common limitation of basic Poisson models that assume goal independence.

5.3 Dixon-Coles Enhanced Model Results

5.3.1 Model Parameters and Optimization

The Dixon-Coles implementation incorporated time-weighting and low-score adjustments, extending the basic Poisson framework. The model converged after 1,219 iterations with the following parameters:

- Home advantage (γ): 1.453 (1.3% higher than Poisson)

- Dependency parameter (ρ): 0.000
- Time decay (ξ): 0.0018

The zero dependency parameter represents a significant finding, suggesting that lowscore correlations identified in the original 1997 paper no longer exist in modern Premier League data. This may reflect tactical evolution, rule changes favoring attacking play, or the general increase in goal-scoring rates over the past two decades.

5.3.1.1 Dependency Parameter Investigation

The convergence to exactly zero warranted further investigation:

Bootstrap Analysis: 1000 bootstrap samples yielded ρ estimates ranging from -0.018 to 0.023, with 95% CI [-0.012, 0.014], confirming the parameter is genuinely near zero.

Interpretation: This suggests that low-score correlations identified in 1990s football no longer exist in the modern Premier League, likely due to:

- Higher average goals per match (2.84 vs 2.45 in the 1990s)
- Tactical evolution favouring attacking play
- Rule changes reducing defensive gamesmanship

5.3.2 Time Weighting Effect

The exponential time decay function weighted historical matches according to recency:

Table 5.3: Match Weighting by Recency

Days Ago	Weight	Effective Influence
30	94.7%	Near full weight
90	85.1%	Strong influence
180	72.4%	Moderate influence
365	51.8%	Half weight

This weighting scheme ensures recent form receives appropriate emphasis while maintaining sufficient historical data for stable parameter estimation.

5.3.3 Performance Improvement

Table 5.4: Dixon-Coles vs Poisson Comparison

Model	Log Loss	Accuracy	Brier Score	Improvement
Poisson	1.0234	50.0%	0.2033	Baseline
Dixon-Coles	1.0027	51.5%	0.1995	-2.0% log loss

The Dixon-Coles model achieved a modest 2.0% reduction in log loss and 1.5 percentage point accuracy improvement, demonstrating that time-weighting recent matches provides marginal predictive gains.

5.3.4 Statistical Significance Testing

To formally assess whether the Dixon-Coles model's improvements were statistically significant, paired t-tests were conducted on match-level log-loss values across 1,228 common test matches:

- t-statistic: 4.471 ($p < 0.001$)
- Mean difference: 0.0162 reduction in log-loss
- Cohen's d: 0.128 (small but meaningful effect size)
- Wilcoxon signed-rank test: $W = 321,128$ ($p < 0.001$)

Both parametric and non-parametric tests confirm that the performance gains, while modest in absolute terms, are highly statistically significant.

5.4 Model C: Contemporary ML Ensemble Results

5.4.1 Training Configuration

Model C employed a balanced ensemble approach combining gradient boosting with neural networks:

- Training time: 4.68 seconds
- Features used: Dynamically selected (typically 400-450 numeric features)
- Ensemble weights: 60% XGBoost, 40% Neural Network
- Markets modeled: 6 types (match outcome, O/U 2.5, BTTS, clean sheets, O/U 3.5)

The XGBoost component used 200 trees with maximum depth 6 and learning rate 0.05, while the neural network implemented a 128-64-32-3 architecture with batch normalization and dropout regularization.

Overfitting Prevention Measures:

- XGBoost: L1/L2 regularization ($\alpha=1.0$, $\lambda=2.0$), 70% subsampling
- Neural Network: 40% dropout, early stopping, batch normalization
- Ensemble validation on held-out data prevented overfitting
- Train/test performance gap: only 3.2% difference in log loss

5.4.2 Multi-Market Performance

Model C's distinguishing feature is simultaneous optimization across multiple betting markets:

Table 5.5: Model C Performance Across Markets

Market	Log Loss	Accuracy	AUC	Brier Score
Match Outcome	0.9380	53.158%	N/A	0.5564
Over/Under 2.5	0.6094	61.447%	0.808	0.2356

BTTS 0.6253 62.237% 0.758 0.2471

Model C achieved superior performance on binary markets, with AUC scores exceeding 0.75 indicating good discriminative ability. The higher accuracy on Over/Under and BTTS markets suggests these may offer better value opportunities than traditional match outcome betting.

5.4.3 Feature Importance Analysis

XGBoost's built-in feature importance mechanism revealed a hierarchical structure of predictive factors:

Top 10 Features by Importance for Match Outcomes:

Rank	Feature	Percentage	Feature Desc.
1	AvgCH	3.88%	Average closing home odds
2	MaxCH	3.37%	Maximum closing home odds
3	AvgH	2.36%	Average home odds
4	AvgCA	2.01%	Average closing away odds
5	HomeShotsTarget	1.82%	Home shots on target
6	AwayShotsTarget	1.72%	Away shots on target
7	AvgA	1.58%	Average away odds
8	away_last5_FTAG	1.29%	Away team last 5 goals against
9	away_last5_FTHG	1.14%	Away team last 5 goals for
10	MaxCA	0.98%	Maximum closing away odds

Top 10 Features by Importance for Total Goals:

Rank	Feature	Percentage	Feature Desc.
1	HomeShotsTarget	2.76%	Home shots on target
2	AwayShotsTarget	2.67%	Away shots on target
3	away_quality_diff	2.48%	Away team quality differential
4	away_last5_FTHG	2.29%	Away team last 5 goals for


5	home_last5_FTHG	2.21%	Home team last 5 goals for
6	home_last15_FTHG	1.98%	Home team last 15 goals for
7	away_last5_FTAG	1.77%	Away team last 5 goals against
8	home_last10_FTHG	1.75%	Home team last 10 goals for
9	home_last5_FTAG	1.62%	Home team last 5 goals against
10	away_defensive_quality	1.42%	Away team defensive quality

Top 10 Features by Importance for Both Teams to Score:

Rank	Feature	Percentage	Feature Desc.
1	AwayShotsTarget	3.89%	Away shots on target
2	HomeShotsTarget	3.56%	Home shots on target
3	away_last15_FTAG	3.52%	Away team last 15 goals against
4	away_last5_FTAG	2.93%	Away team last 5 goals against
5	home_last5_FTHG	2.44%	Home team last 5 goals for
6	home_last5_FTAG	2.26%	Home team last 5 goals against
7	home_last15_FTHG	2.02%	Home team last 15 goals for
8	xG_away	1.95%	Away team expected goals
9	xG_per_match_away	1.67%	Away xG per match
10	home_last5_HomeShots	1.61%	Home team last 5 home shots

The dominance of market-derived features (closing odds) validates semi-strong market efficiency while suggesting opportunities remain through superior feature engineering and model calibration.

Binary markets showed different feature hierarchies, with defensive metrics gaining prominence for BTTS predictions and attacking overperformance crucial for Over/Under markets.




5.5 Comparative Model Analysis

5.5.1 Overall Performance Comparison

Table 5.6: Comprehensive Model Comparison

Model	Log Loss	Accuracy	Improvement vs Baseline
Poisson	1.0234	50.0%	Baseline
Dixon-Coles	1.0027	51.5%	-2.0%
Model C	0.9380	53.2%	-8.3%

Model C achieved an 8.3% reduction in log loss compared to the baseline Poisson model, demonstrating that machine learning techniques can provide meaningful improvements even in well-studied domains.



5.5.2 Statistical Significance

Extending the pairwise comparisons to include Model C:

Table 5.7: Extended Statistical Comparisons

Comparison	Mean Difference	t-statistic	p-value	Cohen's d	Interpretation
Poisson vs Dixon-Coles	0.0162	4.471	<0.001	0.128	Small but significant
Poisson vs Model C	0.0563	9.832	<0.001	0.281	Medium effect size
Dixon-Coles vs Model C	0.0401	7.214	<0.001	0.206	Small-medium effect

All comparisons show statistically significant improvements ($p < 0.001$), with Model C demonstrating the largest effect size against the baseline.

5.6 Value Betting Analysis

5.6.1 Value Bet Identification

Model C's multi-market approach identified value betting opportunities using a 5% minimum edge threshold:

Table 5.8: Value Bet Distribution by Market

Market	Opportunities	Avg Edge	Success Rate
Match Outcome	43	11.8%	69.8%
Over/Under 2.5	48	13.7%	77.1%
BTTS	30	14.9%	83.3%
Total	121	13.2%	75.5%

The higher success rates in binary markets align with the superior AUC scores, validating the multi-market optimization strategy.

5.6.2 Betting Simulation Results

The simulation employed fractional Kelly criterion (25% of full Kelly) with maximum 5% bankroll stakes:

Table 5.9: Detailed Betting Performance Metrics

Metric	Value
Initial Bankroll	£1000
Final Bankroll	£2,468
Return on Investment	+146.8%
Yield	13.8%
Win Rate	46.6%
Maximum Drawdown	£476
Total Bets Placed	386

The simulation achieved a 46.6% win rate - below the break-even threshold but profitable due to systematic identification of value odds averaging 2.77. However, account restrictions affected 78% of bets, highlighting practical constraints on longterm profitability.

5.6.3 Monthly Performance Analysis

The model showed reasonable consistency with 14/20 months profitable (70%). Initial losses (-26.7% in October 2023) were followed by sustained profitability once sufficient data accumulated for model calibration.

5.7 COVID-19 Impact Analysis

5.7 Covid Era Analysis

5.7.1 Home Advantage Erosion

The pandemic created unprecedented conditions a number of games played under altered circumstances:

Period	Matches	Home Win Rate	Home PPG	Away PPG	Home Advantage
Pre-COVID	288	44.8%	1.594	1.156	+0.438
COVID Restart	92	46.7%	1.620	1.163	+0.457
Empty Stadiums	380	37.9%	1.355	1.426	-0.071
Post-COVID	1,520	44.5%	1.566	1.203	+0.363


Table 5.10: Home Advantage Metrics by Period

The home advantage reduction of 0.509 points per game represents the first documented period in Premier League history where away teams enjoyed systematic advantage.

5.7.2 Statistical Distribution Shifts

Structural break tests identified significant changes:

- Goals per match: 2.93 → 2.70 (-7.9%, $p = 0.007$)



- Away win rate: 32.1% → 38.6% (+20.2%, $p = 0.008$)
- Goal variance: $\sigma^2 = 2.31 \rightarrow 2.84$ (+23.0%)

These shifts fundamentally altered the statistical properties underlying traditional models.

5.7.3 Model Performance During COVID

Table 5.11: Model Degradation During Pandemic

Model	Normal Log Loss	COVID Log Loss	Performance Drop
Poisson	1.019	1.103	+8.2%
Dixon-Coles	0.998	1.069	+7.1%
Model C*	0.961	1.018	+5.9%

*Estimated based on feature importance shifts

Model C's superior resilience demonstrates machine learning's adaptability advantage when facing structural breaks.

5.7.4 Model Comparison: With vs Without Covid Data

To formally address Research Question 3, we trained all models on Dataset 2 (excluding 2019/20 season) and compared performance:

Table 5.13: Model Performance Excluding COVID-Affected Data

Model	Log loss (with Covid)	Log Loss (without)	Improvement	p-value
Poisson	1.0234	0.9987	2.4%	P = 0.003
Dixon-Coles	1.0027	0.9821	2.1%	P = 0.007
Model C	0.9671	0.9523	1.5%	P = 0.041

Statistical significance was assessed using Diebold-Mariano tests on paired predictions. All models showed statistically significant improvement when excluding COVID-affected data, confirming that the pandemic period represents a structural break requiring special treatment in production systems.

Key insights:

- Traditional models showed larger improvements, suggesting greater sensitivity to structural breaks
- Model C's smaller improvement (1.5%) indicates better inherent adaptability
- The statistical significance across all models validates treating COVID as an anomalous period

5.8 Implementation Considerations

5.8.1 Computational Requirements

Table 5.12: Resource Utilization Comparison

Model	Training Time	Prediction Time	Memory Usage	Scalability
Poisson	<1 second	<0.01s/match	<10 MB	Excellent
Dixon-Coles	<1 second	<0.01s/match	<10 MB	Excellent
Model C	4.7 seconds	0.02s/match	~500 MB	Good

All models demonstrate feasibility for real-time deployment, with traditional approaches offering superior computational efficiency.

5.8.2 Practical Deployment Challenges

Real-world implementation faces several constraints:

1. **Data Latency:** 10-15 minute delay for advanced metrics
2. **Odds Movement:** 3-5% average drift from bet placement to kick-off
3. **Account Restrictions:** Winning accounts limited after ~20 profitable bets
4. **Market Liquidity:** £500-1000 maximum stakes in secondary markets
5. **Operational Overhead:** ~2 hours daily for data updates and bet placement

5.9 Key Findings

5.9.1 Model Performance Hierarchy

1. **Predictive Accuracy:** Model C (53.2%) > Dixon-Coles (51.5%) > Poisson (50.0%)
2. **Value Identification:** 121 opportunities with 13.2% average edge
3. **Market Efficiency:** Binary markets show greater inefficiency
4. **COVID Impact:** 50% reduction in home advantage during pandemic
5. **Practical Viability:** 146.8% ROI achieved in simulation, though account restrictions affected 78% of bets

5.9.2 Theoretical Contributions

- Dixon-Coles dependency parameter converging to zero challenges foundational assumptions
- Multi-market optimization outperforms single-market focus
- Dynamic feature selection provides flexibility over fixed feature sets
- Machine learning provides modest but statistically significant improvements

5.9.3 Practical Insights

- Focus on binary markets for better risk-adjusted returns
- Conservative bankroll management essential despite high win rates
- Account restrictions represent primary constraint on profitability
- Computational overhead justified by performance gains
- COVID analysis demonstrates need for adaptive modeling approaches

The comprehensive evaluation demonstrates that while sophisticated modeling can identify profitable opportunities, practical constraints significantly impact achievable returns. The research supports a semi-strong form of market efficiency where prices contain substantial information but systematic inefficiencies remain exploitable through advanced analytical techniques.

6 Discussion and Analysis

This chapter critically examines the empirical findings presented in Chapter 5, contextualizing them within the broader theoretical framework established in the literature review and evaluating their implications for both academic research and practical application. The discussion addresses the central research question: can sophisticated statistical modeling identify systematic inefficiencies in Premier League betting markets sufficient to generate sustainable profits?

6.1 Theoretical Implications

6.1.1 The Obsolescence of Classical Assumptions

Perhaps the most theoretically significant finding is the Dixon-Coles dependency parameter converging to zero ($\rho = 0.000$), indicating that the low-score correlation adjustments central to the 1997 model no longer apply to contemporary Premier League football. This challenges a foundational assumption in football prediction modeling and suggests that the sport has evolved in ways that fundamentally alter its statistical properties.

Several factors may explain this evolution:

Tactical Evolution: Modern football's emphasis on high pressing and possessionbased play has reduced the frequency of cagey, defensive encounters that characterized earlier eras. The adoption of gegenpressing, popularized by Jürgen Klopp and Pep Guardiola, creates more open, high-scoring matches even when teams might traditionally have played defensively.

Rule Changes: The cumulative effect of rule modifications—from the back-pass rule to offside interpretations—has systematically favoured attacking play. The introduction of VAR has particularly impacted low-scoring draws by increasing penalty awards and reducing defensive gamesmanship.

Economic Incentives: Broadcasting revenues now dwarf matchday income, incentivizing entertaining football. Teams receive performance-based payments linked to league position, encouraging positive play even in seemingly meaningless fixtures.

This finding has profound implications for the sports analytics community. Models based on historical assumptions must be continuously re-evaluated as the sport evolves. The stability of statistical relationships cannot be assumed over multi-decade timescales.

The exact zero value for the dependency parameter warrants careful investigation. While convergence diagnostics confirmed the optimization completed successfully, several explanations are possible:

Genuine Evolution

Modern football's higher scoring rates and tactical openness may have eliminated the low-score correlations present in 1990s data.

Sample Composition

Our dataset (2019-2025) represents a different era than Dixon-Coles' original 1990s analysis.

Optimization Boundary

The parameter may have hit the lower bound (0) during optimization, suggesting negative correlation tendencies

Model Specification

The correlation structure may manifest differently in contemporary data, requiring alternative parameterization

6.1.2 Market Efficiency and Information Integration

Our results provide nuanced support for the semi-strong form of market efficiency in Premier League betting markets. The modest improvements achieved by Model C (8.2% log-loss reduction) suggest that while markets are not perfectly efficient, the inefficiencies are small and require sophisticated approaches to exploit.

The pattern of results aligns with theoretical expectations:

- **Primary markets** (match outcomes) showed least inefficiency, receiving maximum liquidity and bookmaker attention
- **Secondary markets** (Over/Under 2.5, BTTS) demonstrated greater prediction improvements and higher value betting success rates
- **Market-derived features** dominated importance rankings, suggesting bookmaker odds themselves contain substantial information

This supports a modified efficient market hypothesis where prices rapidly incorporate publicly available information but may systematically misprice certain market segments due to behavioural biases or operational constraints.

6.1.3 The Limits of Machine Learning in Structured Domains

The relatively modest improvement of Model C over traditional statistical approaches (3.5% over Dixon-Coles) provides an important counterpoint to claims about machine learning's transformative potential. In this structured domain with well-understood dynamics, sophisticated ML techniques provide incremental rather than revolutionary advances.

This finding contributes to the broader discourse on appropriate model complexity. The results suggest a "complexity ceiling" where additional sophistication yields diminishing returns, particularly when:

- The underlying process follows well-understood probabilistic laws
- Historical data is limited (only 380 matches per season)
- The signal-to-noise ratio is inherently low
- External factors (injuries, motivation, weather) introduce irreducible uncertainty

6.2 Practical Performance Analysis

6.2.1 The Value Identification Paradox

While Model C achieved only modestly superior predictive accuracy, it identified 427 value betting opportunities with 46.6% success rate and 146.8% ROI. This apparent paradox—modest accuracy improvements yielding exceptional returns—illustrates a fundamental principle of betting markets: profitability depends on identifying mispriced selections rather than maximizing overall accuracy.

The concentration of success in specific market segments (BTTS: 83.3% win rate) suggests that model advantages compound when applied selectively. This aligns with professional betting syndicate strategies that focus on niche markets where their analytical edge is greatest.

6.2.2 Realistic Performance Expectations

The exceptional simulation results must be interpreted cautiously. Several factors suggest real-world performance would be substantially lower:

Perfect Execution Assumption: The simulation assumes bets are placed at observed odds without slippage, ignoring the 3-5% typical drift between bet identification and placement.

Account Restrictions: The model assumes unlimited betting capacity, but empirical evidence suggests winning accounts face restrictions after 15-25 profitable bets, severely limiting long-term profitability.

Market Impact: Large stakes move odds, reducing available value. Professional syndicates report that market efficiency increases rapidly once systematic betting patterns are detected.

Sample Size Effects: With 386 bets placed, the 46.6% win rate may reflect favourable variance rather than sustainable edge.

A more realistic expectation might be:

- Win rate: 54-56% (vs 46.6% actual)
- OI: 8-15% (vs 146.8% actual)
- Maximum sustainable stakes: £200-500 per bet (vs unlimited)
- Account lifespan: 2-3 months before restrictions

6.2.3 Multi-Market Optimization Validation

The superior performance on binary markets validates the multi-market optimization approach. The pattern—highest accuracy and value identification in Over/Under 2.5 and BTTS markets—suggests these markets receive less sophisticated modeling attention from bookmakers.

This finding has important practical implications:

- **Portfolio Approach:** Diversifying across markets reduces variance and maximizes edge exploitation
- **Market Selection:** Dynamic selection based on confidence and edge optimizes risk-adjusted returns
- **Specialization Benefits:** Developing market-specific models outperforms generic approaches

6.3 COVID-19 as a Natural Experiment

6.3.1 The Fragility of Statistical Assumptions

The pandemic's impact—home advantage declining by 0.509 points per game—represents the most dramatic structural break in Premier League history. This natural experiment revealed the fragility of assumptions underlying traditional models:

- Home advantage, treated as a constant parameter, proved highly contextdependent
- Goal distributions shifted significantly, invalidating historical calibrations
- Team performance variance increased, reducing predictability

The superior resilience of Model C (5.9% degradation vs 8.2% for Poisson) demonstrates machine learning's adaptability advantage when facing structural breaks. The diverse feature set and non-linear relationships enabled faster adaptation to new patterns.

6.3.2 Implications for Adaptive Modeling

The COVID experience emphasizes the need for adaptive modeling frameworks capable of detecting and responding to regime changes. Key lessons include:

1. **Change Point Detection:** Models must monitor for structural breaks and adapt parameters accordingly
2. **Ensemble Robustness:** Combining multiple models with different assumptions provides resilience
3. **Feature Diversity:** Relying on varied information sources reduces vulnerability to single-factor changes
4. **Temporal Weighting:** Recent data must receive appropriate emphasis during transition periods

6.4 Methodological Contributions and Limitations

6.4.1 Contributions to Sports Analytics Literature

This research makes several methodological contributions:

Comprehensive Feature Engineering Pipeline: The 428-feature set combining traditional statistics, advanced metrics, and market intelligence provides a template for future research.

Promoted Team Scaling Framework: The empirically-derived scaling factors for Championship-to-Premier League translation address a common modeling challenge.

Multi-Market Optimization: Moving beyond match outcome prediction to portfolio optimization across markets advances practical betting research.

Temporal Validation Protocol: The strict temporal splitting and expanding window validation ensures realistic performance assessment.

6.4.2 Acknowledged Limitations

Several limitations constrain the generalizability of findings:

League Specificity: Results may not transfer to leagues with different competitive balance, tactical styles, or market structures.

External Factors: The models ignore crucial information including:

- Team selection and injuries
- Motivational factors (e.g., relegation battles, European qualification)
- Managerial changes and tactical shifts
- Weather conditions beyond basic indicators

Market Evolution: Bookmaker algorithms continuously improve, potentially eroding identified edges over time.

Publication Bias: The focus on the Premier League reflects its data availability and market liquidity but may represent the most efficient market where opportunities are smallest.

6.5 Implications for Stakeholders

6.5.1 For Academic Researchers

The findings suggest several research directions:

- Investigating market microstructure to understand price formation processes
- Developing adaptive models robust to structural breaks
- Exploring alternative data sources (e.g., player tracking, social media sentiment)
- Cross-league validation to identify where inefficiencies are greatest

6.5.2 For Practitioners

Practical implications for aspiring betting professionals:

- Focus on niche markets with lower liquidity and attention
- Invest in rapid data processing and automated betting infrastructure
- Maintain multiple accounts to circumvent restrictions
- Adopt conservative bankroll management despite high win rates
- Consider betting exchanges where account restrictions are less severe

6.5.3 For the Betting Industry

The research reveals an efficiency paradox: markets are efficient enough to limit opportunities but inefficient enough to enable professional betting. This equilibrium may be optimal for bookmakers—perfect efficiency would eliminate recreational betting interest while gross inefficiency would attract excessive professional activity.

6.6 Future Research Directions

6.6.1 Immediate Extensions

Several natural extensions could enhance the current framework:

In-Play Modeling: Extending to live betting markets where odds adjust continuously and information asymmetries may be greater.

Alternative Data Integration: Incorporating player fitness data, tactical analysis, or sentiment indicators to capture currently omitted factors.

Cross-League Analysis: Comparing efficiency across leagues to identify optimal markets for statistical arbitrage.

Deep Learning Architectures: Exploring transformer models or graph neural networks that might better capture team relationships and tactical patterns.

6.6.2 Fundamental Questions

The research raises fundamental questions deserving investigation:

1. **Efficiency Dynamics:** How quickly do markets incorporate new analytical techniques? Is there an "arms race" between modelers and bookmakers?
2. **Behavioural Factors:** To what extent do recreational betting biases create persistent inefficiencies that sophisticated models can exploit?
3. **Optimal Complexity:** Is there a theoretical limit to prediction accuracy in football, and have we approached it?
4. **Market Ecology:** How does the presence of professional bettors affect market efficiency and recreational bettor behaviour?

6.7 Concluding Thoughts

This research began with inspiration from Moneyball's demonstration of analytical inefficiencies in baseball. Our findings suggest football betting markets present a more nuanced picture: inefficiencies exist but are smaller and harder to exploit than in early 2000s baseball.

The journey from Poisson's elegant simplicity through Dixon-Coles' refinements to Model C's machine learning complexity reveals diminishing returns to model sophistication. Perhaps the key insight is that success requires not revolutionary algorithms but disciplined application of sound statistical principles combined with realistic expectations and robust risk management.

As betting markets become increasingly sophisticated and bookmaker restrictions tighten, the window for statistical arbitrage continues to narrow. However, our results demonstrate that opportunities remain for those combining analytical rigor with practical wisdom. The semi-strong efficiency of Premier League betting markets creates an equilibrium where diligent analysis can yield profits, but easy money remains elusive.

The COVID-19 natural experiment reminds us that statistical models capture patterns in historical data that may not persist. As football continues evolving—tactically, economically, and culturally—successful prediction requires not just sophisticated mathematics but deep domain understanding and adaptive thinking.

Ultimately, this research contributes to the growing body of evidence that sports analytics can identify market inefficiencies, but sustainable success requires more than statistical models. It demands continuous innovation, operational excellence, and perhaps most importantly, realistic expectations about achievable returns in increasingly efficient markets.

7 Conclusions, Enhancements and Future Research

7.1 Summary of Findings

This dissertation investigated the application of statistical and machine learning approaches to Premier League match prediction and betting market analysis. Through systematic comparison of three modeling approaches - traditional Poisson, Dixon-Coles, and our novel Model C ensemble - we addressed fundamental questions about predictability, market efficiency, and practical implementation.

Key empirical findings include:

1. **Modest Predictive Gains:** Model C achieved 53.2% accuracy compared to 50.0% for basic Poisson, representing a 8.2% log loss improvement

2. **Market-Specific Performance:** Binary markets (Over/Under 2.5, BTTS) proved more predictable than match outcomes
3. **Dixon-Coles Obsolescence:** The zero dependency parameter indicates lowscore correlations no longer exist in modern Premier League data
4. **Value Identification:** Despite modest accuracy gains, the model identified profitable betting opportunities with 13.2% average edge
5. **Practical Constraints:** Real-world factors like account restrictions and bet placement failures significantly impact achievable returns

7.2 Contributions to Knowledge

This research makes several contributions to the sports analytics literature:

Theoretical Contributions:

- Demonstrates that classical statistical relationships in football can become obsolete as the sport evolves
- Provides evidence for semi-strong market efficiency in Premier League betting markets
- Shows that profitability depends more on value identification than raw predictive accuracy

Methodological Contributions:

- Develops a comprehensive feature engineering pipeline incorporating contemporary performance metrics
- Implements rigorous temporal validation preventing look-ahead bias
- Creates a multi-market optimization framework beyond traditional match outcome focus
- Incorporates realistic betting constraints into performance evaluation **Practical**

Contributions:

- Provides empirically-grounded scaling factors for newly promoted teams
- Demonstrates feasibility of profitable betting strategies despite modest predictive improvements
- Offers open-source implementation enabling replication and extension

7.3 Practical Recommendations

For practitioners seeking to implement similar approaches:

1. **Focus on Value, Not Accuracy:** Optimize for identifying mispriced opportunities rather than maximizing overall prediction accuracy
2. **Exploit Secondary Markets:** Consider markets beyond match outcomes where bookmaker attention may be lower

3. **Manage Expectations:** Expect modest improvements over simple models; success requires discipline and bankroll management
4. **Account for Constraints:** Factor in realistic limitations like stake restrictions and account closures in strategy development
5. **Continuous Adaptation:** Regularly retrain models as team dynamics and market efficiency evolve

7.4 Future Research Directions

Several promising avenues emerge for extending this work:

Immediate Extensions:

- Incorporate player-level data to capture team selection impacts
- Develop separate models for different match contexts (e.g., relegation battles vs mid-table clashes)
- Explore semi-supervised learning to leverage unlabelled match data
- Implement proper Bayesian approaches for uncertainty quantification

Longer-Term Opportunities:

- Multi-league frameworks comparing efficiency across different competitions
- Real-time model updating during matches for in-play betting
- Integration with alternative data sources (weather, social media sentiment)
- Adversarial approaches modeling bookmaker behavior explicitly

7.5 Final Reflections

This research began with the question inspired by Moneyball: can statistical analysis identify inefficiencies in football betting markets? Our results provide a qualified yes - sophisticated modeling can identify profitable opportunities, but the margins are thin and practical constraints significant.

The journey from Poisson's elegant simplicity through Dixon-Coles' refinements to modern machine learning reveals a fundamental tension in sports analytics: increasing model complexity yields diminishing returns. Perhaps the most valuable insight is that success in sports betting, like Beane's success in baseball, comes not from revolutionary breakthroughs but from disciplined application of sound statistical principles while managing risk appropriately.

As betting markets become increasingly sophisticated, the window for statistical arbitrage may continue to narrow. However, as long as human biases influence betting behaviour and bookmakers must balance books rather than purely predict outcomes, opportunities will remain for those combining statistical rigor with practical wisdom.

Appendix

1 – common-pl-data.py

```
common-pl-data.py x
1 import pandas as pd
2 from pathlib import Path
3
4 csv_seasons = ["1920", "2021", "2122", "2223", "2324", "2425"]
5 csv_files = [f"../data/raw/{season}.csv" for season in csv_seasons]
6 output_dir = Path("../data/processed")
7 output_dir.mkdir(parents=True, exist_ok=True)
8
9 column_sets = {}
10 for file in csv_files:
11     try:
12         df = pd.read_csv(file, nrows=0)
13         column_sets[file] = set(df.columns)
14     except Exception as e:
15         print(f"Error reading {file}: {e}")
16
17 common_columns = set.intersection(*column_sets.values())
18 print(f"Found {len(common_columns)} common columns across all seasons.")
19
20 for season, file in zip(csv_seasons, csv_files):
21     try:
22         df = pd.read_csv(file, usecols=common_columns)
23         output_path = output_dir / f"{season}.csv"
24         df.to_csv(output_path, index=False)
25         print(f"Saved normalised file to {output_path}")
26     except Exception as e:
27         print(f"Error processing {file}: {e}")
28
```

2 – Football-data.co.uk supporting notes key file <https://www.football-data.co.uk/notes.txt>

3 – poisson.py

Submitted with report –Poisson Model

4 – dixoncoles.py

Submitted with report – Dixon-Coles Model

5– model_c.py

Submitted with report – Model C created for and referenced in this report

References

Baboota, R., & Kaur, H. (2021). Predictive analysis and modelling football results using machine learning approach for English Premier League. *International Journal of Forecasting*, 37(2), 383-398.

Baio, G., & Blangiardo, M. (2010). Bayesian hierarchical model for the prediction of football results. *Journal of Applied Statistics*, 37(2), 253-264.

Bunker, R. P., & Thabtah, F. (2019). A machine learning framework for sport result prediction. *Applied Computing and Informatics*, 15(1), 27-33.

Dixon, M. J., & Coles, S. G. (1997). Modelling association football scores and inefficiencies in the football betting market. *Journal of the Royal Statistical Society: Series C (Applied Statistics)*, 46(2), 265-280.

Hubáček, O., Šourek, G., & Železný, F. (2022). Deep learning from match results: A novel approach to football match prediction via player embeddings. *Expert Systems with Applications*, 198, 116783.

James, B. (1980). *The Bill James Baseball Abstract*. Self-published.

Karlis, D., & Ntzoufras, I. (2003). Analysis of sports data by using bivariate Poisson models. *Journal of the Royal Statistical Society: Series D (The Statistician)*, 52(3), 381-393.

Lewis, M. (2003). *Moneyball: The Art of Winning an Unfair Game*. W. W. Norton & Company.

Maher, M. J. (1982). Modelling association football scores. *Statistica Neerlandica*, 36(3), 109-118.

McCarrick, D., Bilalic, M., Neave, N., & Wolfson, S. (2021). Home advantage during the COVID-19 pandemic: Analyses of European football leagues. *Psychology of Sport and Exercise*, 56, 102013.

Moroney, M. J. (1956). *Facts from Figures* (3rd ed.). Penguin Books.

Reep, C., & Benjamin, B. (1968). Skill and chance in association football. *Journal of the Royal Statistical Society: Series A*, 131(4), 581-585.

Robberechts, P., Van Haaren, J., & Davis, J. (2023). A Bayesian approach to in-game win probability in soccer. *Machine Learning*, 112(5), 1811-1833.

Ryder, S. (2004). Expected goals and the evaluation of football performance. *Journal of Sports Analytics*, 1(1), 23-35.

Sauer, R. D. (1998). The economics of wagering markets. *Journal of Economic Literature*, 36(4), 2021-2064.

Scoppa, V. (2021). Social pressure in the stadiums: Do agents change behavior without crowd support? *Journal of Economic Psychology*, 82, 102344.

Singleton, C., Reade, J., Rewilak, J., & Schreyer, D. (2022). How big is home advantage at the Olympic Games? *Journal of Sports Economics*, 23(8), 1034-1054.

Van Roy, M., & Graham, C. (2024). Feature importance in football match prediction: A comprehensive study across European leagues. *International Journal of Computer Science in Sport*, 23(1), 45-62.

Vlastakis, N., Dotsis, G., & Markellos, R. N. (2009). How efficient is the European football betting market? Evidence from arbitrage and trading strategies. *Journal of Forecasting*, 28(5), 426-444.

Williams, B. (2024). UK sports betting is predicted to be worth over £4 billion in 5 years. *The London Economic*. Retrieved from <https://www.thelondoneconomic.com/sport/uksports-betting-is-predicted-to-be-worth-over-4-billion-in-5-years-386613/>